

Exam Notes: The Rothschild–Stiglitz Model

Screening in Competitive Insurance Markets

Information Economics I · HT 2026

1. Setting and Motivation

The Rothschild–Stiglitz (1976) model studies a competitive insurance market with **adverse selection**: insurers cannot observe the riskiness of individual customers. It is the canonical model of **screening** — the uninformed side (insurers) designs contracts to induce the informed side (customers) to reveal their type through self-selection.

Key question

Can competitive markets function efficiently when buyers know more about their own risk than sellers?

Context in the lectures. The model is an instance of a principal–agent problem with *hidden information* where information is asymmetric *before* contracting, and the uninformed party (insurer/principal) moves first by offering a menu of contracts (screening). Compare with:

- **Signalling** (Spence): informed party moves first.
- **Monopolistic screening** (lecture slides): a single principal designs the menu. R–S studies the *competitive* analogue.

2. Model Primitives

2.1 Individuals

- Initial wealth W ; potential loss (accident) of size d .
- Two states of nature: *no accident* (income W) and *accident* (income $W - d$).
- An insurance contract $\alpha = (\alpha_1, \alpha_2)$: the individual pays premium α_1 and receives net benefit α_2 if an accident occurs.
- With insurance: income is $(W - \alpha_1, W - d + \alpha_2)$ where $\alpha_2 = \hat{\alpha}_2 - \alpha_1$.

Preferences. Expected utility with a von Neumann–Morgenstern utility function:

$$V(p, W_1, W_2) = (1 - p)U(W_1) + pU(W_2), \quad U'' < 0 \text{ (risk averse)}.$$

Two types of customer:

- **Low-risk:** accident probability p^L .
- **High-risk:** accident probability $p^H > p^L$.
- Fraction of low-risk: λ . Average probability: $\bar{p} = \lambda p^L + (1 - \lambda)p^H$.

Information asymmetry

Individuals know their own type (p^L or p^H). Insurance companies **cannot** observe individual types — they only know the population distribution $(\lambda, 1 - \lambda)$.

2.2 Insurance Companies

- Risk-neutral: care only about expected profits.
- Expected profit from selling contract α to a type- p individual:

$$\pi(p, \alpha) = (1 - p) \alpha_1 - p \alpha_2 = \alpha_1 - p(\alpha_1 + \alpha_2).$$

- Free entry $\Rightarrow$ any contract that earns positive expected profits will be offered.
- Each contract must individually break even (zero expected profit) in equilibrium.

2.3 Contract Space and Graphical Representation

The state-space diagram has W_1 (income if no accident) on one axis and W_2 (income if accident) on the other. The endowment point $E = (W, W - d)$ lies below the 45° line.

- The **fair-odds line** for type p passes through E with slope $(1 - p)/p$.
- The 45° line represents **full insurance** ($W_1 = W_2$).
- Indifference curves: slope $= \frac{(1-p) U'(W_1)}{p U'(W_2)}$. On the 45° line, slope $= (1 - p)/p$, independent of U .

Single-crossing property

At any point in the (W_1, W_2) space, the indifference curve of the *low-risk* type is **steeper** than that of the high-risk type:

$$\text{slope of } U^L = \frac{(1 - p^L) U'(W_1)}{p^L U'(W_2)} > \frac{(1 - p^H) U'(W_1)}{p^H U'(W_2)} = \text{slope of } U^H.$$

This is crucial for the entire screening argument.




Figure 1: State-space diagram. EL and EH are the zero-profit (fair-odds) lines for low- and high-risk types. EF is the market-average fair-odds line. Under symmetric information, each type gets full insurance at the 45° line on their respective fair-odds line.

3. Equilibrium Concept

Rothschild–Stiglitz (Nash) equilibrium

An equilibrium is a set of contracts $\{\alpha^*\}$ such that, when customers choose to maximise expected utility:

1. No contract in the equilibrium set makes negative expected profits.
2. There is **no contract outside** the equilibrium set that, if offered, would make a non-negative expected profit.

This is a **Cournot–Nash** concept: each firm takes others' contracts as given.

Two possible equilibrium types:

- **Pooling equilibrium:** both types buy the same contract.
- **Separating equilibrium:** different types buy different contracts.

4. Benchmark: Symmetric Information (Full Information)

If insurers can observe types, the market splits into two sub-markets. Each type gets:

- **Full insurance** at actuarially fair odds: the contract lies at the intersection of the type's fair-odds line and the 45° line.
- High-risk types pay premium $p^H d$; low-risk types pay $p^L d$.
- Both achieve **complete insurance**: $W_1 = W_2 = W - p^i d$.
- Zero profits for firms.

This is the first-best outcome: efficient risk-sharing with no distortions.

5. Asymmetric Information: No Pooling Equilibrium

Theorem: No pooling equilibrium exists

There cannot be a pooling equilibrium in the R–S model.

Proof sketch:

1. Suppose α is a pooling equilibrium bought by both types.
2. Zero-profit condition $\Rightarrow \alpha$ lies on the *market* fair-odds line EF (slope $(1 - \bar{p})/\bar{p}$).
3. By the **single-crossing property**, the low-risk indifference curve through α is steeper than the high-risk curve at α .
4. Therefore there exists a contract β near α that:
 - Low-risk types *prefer* to α (it is above $\bar{U}^L$);
 - High-risk types *do not* prefer to α (it is below $\bar{U}^H$).
5. Since only low-risk types buy β , expected profit $\pi(p^L, \beta) > 0$.
6. This contradicts condition (ii) of equilibrium. □

Why does steepness create a cream-skimming opportunity? The key geometric insight is shown in Figure 2. Because the two indifference curves cross at α but have *different slopes*, they must diverge on either side. To the *right* of α (the direction of less insurance coverage), the steeper low-risk curve $\bar{U}^L$ drops below the flatter high-risk curve $\bar{U}^H$, opening a **wedge-shaped gap**. Any contract β placed inside this wedge is simultaneously:

- **Above** $\bar{U}^L$: so low-risk types prefer β to α .
- **Below** $\bar{U}^H$: so high-risk types prefer α to β — they stay put.

Since β attracts *only* low-risk types (who have $p^L < \bar{p}$), an insurer offering β makes a strictly positive expected profit.


Figure 2: The cream-skimming wedge — why no pooling equilibrium exists

Figure 2: **Why no pooling equilibrium exists.** At the pooling contract α , the single-crossing property means $\bar{U}^L$ (steeper) and $\bar{U}^H$ (flatter) diverge to the right, opening a green *wedge*. Any contract β in the wedge is preferred by L-types but not H-types. Since only L-types buy β , it earns $\pi(p^L, \beta) > 0$, violating the equilibrium condition.

6. Separating Equilibrium: Characterisation

Unique candidate separating equilibrium

If an equilibrium exists, it must be separating with contracts (α^H, α^L) such that:

1. α^H gives **full insurance** at high-risk fair odds: α^H is at the intersection of EH and the 45° line.
2. α^L gives **partial insurance** at low-risk fair odds: α^L is on line EL at the point where the *high-risk* indifference curve through α^H intersects EL .
3. Both contracts make zero expected profits.

Construction:

1. High-risk contract: the contract on EH most preferred by high-risk types gives full insurance $\Rightarrow \alpha^H$ on 45° line $\cap EH$.
2. Low-risk contract must satisfy the **incentive-compatibility (self-selection) constraint**: high-risk types must weakly prefer α^H to α^L .

$$V(p^H, \alpha^H) \geq V(p^H, \alpha^L)$$

That is, α^L must lie on or below U^H , the high-risk indifference curve through α^H .

3. The best such contract for low-risk types on line EL : the intersection $\alpha^L = EL \cap U^H$.




Figure 3: **The unique candidate separating equilibrium** (α^H, α^L) . High-risk types get full insurance at their fair odds (α^H on $45^\circ \cap EH$). Low-risk types get only *partial* insurance at their fair odds (α^L on $EL \cap U^H$), because the IC constraint forces α^L to lie on the high-risk indifference curve U^H through α^H .

Key features of the separating equilibrium

- **High-risk types:** full insurance at actuarially fair odds — *no distortion at the top* (they get their first-best contract).
- **Low-risk types:** only partial insurance at fair odds — they are *distorted downward* to prevent high-risk mimicking.
- Contracts specify **both price and quantity** (not just a price at which customers can buy any amount).
- This is a **self-selection mechanism**: the menu of contracts induces truthful revelation of type.

7. Non-Existence of Equilibrium

The candidate separating equilibrium (α^H, α^L) may itself be “broken.”

Equilibrium may not exist

If there are relatively **many low-risk** individuals (i.e. λ is large), the separating equilibrium does not survive, and since no pooling equilibrium exists either, **no equilibrium exists**.

Mechanism (see Figure 4):

1. Consider a contract γ that lies above *both* U^H (through α^H) and $\bar{U}^L$ (through α^L). Both types prefer γ to their separating contract.
2. If γ is offered, both types buy it.
3. Profitability depends on the market mix:
 - If λ is **small** (many H) $\Rightarrow$ market odds line EF close to $EH \Rightarrow \gamma$ makes a loss $\Rightarrow$ equilibrium survives.
 - If λ is **large** (many L) $\Rightarrow$ market odds line EF' close to $EL \Rightarrow \gamma$ makes a profit $\Rightarrow$ equilibrium broken.

Figure 4: Non-existence — pooling deviation breaks equilibrium when λ small




Figure 4: Non-existence of equilibrium. If the fraction of low-risk types λ is large, the market-odds line is close to EL , and a pooling contract γ that both types prefer to (α^H, α^L) can earn positive profits. Since (a) the separating equilibrium is broken, and (b) no pooling equilibrium exists (Section 5), **no equilibrium exists**.

Conditions for existence / non-existence

Equilibrium **does not exist** when:

- The cost of pooling is low for the low-risk type (many low-risk individuals, so few high-risk to subsidise, or $p^H - p^L$ small).

- The cost of separating is high (individuals are very risk-averse, so partial insurance is very costly).

Equilibrium **exists** when λ is sufficiently **small** (many high-risk individuals), $p^H - p^L$ is large, or individuals are not too risk-averse.

8. Welfare Properties

Welfare results

1. **Negative externality:** High-risk individuals impose a *dissipative* externality on low-risk individuals. Low-risk types are worse off than under symmetric information, while high-risk types are *no better off* than they would be alone (they get their full-info contract).
2. **Pareto sub-optimality:** The separating equilibrium may not be (constrained) Pareto optimal. There may exist a pair of contracts that break even *together* (cross-subsidisation) and make both types better off.
3. **Value of information:** If high-risk individuals could credibly reveal their type, *everyone* would be better off.

Figure 5: Cross-subsidisation — Pareto improvement, but unstable




Figure 5: **Cross-subsidisation can Pareto-dominate the separating equilibrium.** A pair $(\alpha^{H'}, \alpha^{L'})$ that breaks even *jointly* can make both types better off. But it is not stable under free entry: an entrant could cream-skin the low-risk types from $\alpha^{L'}$, destroying the cross-subsidy.

9. Robustness and Extensions

9.1 Price vs. Price-and-Quantity Competition

R–S assume firms offer contracts specifying *both* a premium and a coverage amount (price-and-quantity competition), rather than just a unit price of insurance. This is justified by:

- **Moral hazard:** firms want to limit total coverage to prevent perverse incentives.
- **Competitive dominance:** price-and-quantity competition dominates pure price competition.
- **Institutional features:** many real insurance contracts specify that they are void if other coverage exists, effectively enforcing exclusivity.

9.2 Alternative Equilibrium Concepts

- **Wilson equilibrium (1977):** No entrant makes a profit after loss-making incumbents withdraw. Always exists.
- **Riley reactive equilibrium:** Firms anticipate competitor reactions. Separating equilibrium always survives.
- **Informationally consistent equilibrium (Spence):** Any break-even allocation given the purchasing pool. Many equilibria exist.

9.3 Continuum of Types

With a continuous distribution of accident probabilities, R–S equilibrium **never** exists (Riley, 1979). There are always types close enough that pooling them is beneficial.

10. Connection to the Lecture Slides

Feature	Monopolistic Screening	R–S (Competition)
No. of principals	One	Many (free entry)
Profit condition	Total profits ≥ 0	Each contract ≥ 0
No distortion at top	$\hat{e}_H = e_H^*$	$\alpha^H = \text{full insurance}$
Distortion at bottom	$\hat{e}_L < e_L^*$	$\alpha^L = \text{partial insurance}$
Info rent to “good” type	$u_H > \bar{u}$	Low-risk: no rent
Cross-subsidisation	Not in equilibrium	Not in equilibrium
Existence guaranteed?	Yes	Not always

The key difference: in competitive screening, the zero-profit condition is per-contract (no cross-subsidy), which can cause non-existence.

11. Exam Practice: Competitive, Social Planner, Monopolist

Exam Question

Consider the Rothschild–Stiglitz model of competitive provision of insurance, where consumers are privately informed of their type, and the fraction of low-risk consumers is λ .

(a) Explain why high-risk consumers get full insurance in the separating equilibrium. What is the condition on λ that ensures equilibrium exists, and why?

(b) Assume λ is such that a social planner can effect a Pareto improvement on the separating equilibrium (without incurring a deficit). What trade-off does she consider? Why do high-risk consumers again get full insurance? What condition on λ ensures a Pareto improvement is possible?

(c) Consider a risk-neutral monopolist. What policies does he offer when λ is low (many high-risk) vs. λ is high (many low-risk)? Will high-risk consumers again get full insurance?

Part (a): Competitive Separating Equilibrium

Why do high-risk consumers receive full insurance?

In the separating equilibrium, each contract must break even individually (zero expected profit under free entry). So α^H lies on the high-risk fair-odds line EH .

The key observation is that *no incentive-compatibility constraint binds on the high-risk type's contract from above*. Nobody wants to mimic a high-risk person — the low-risk would never choose the expensive high-risk policy because they face better odds on their own fair-odds line. Since the only constraint on α^H is zero profit (it lies on EH), we simply choose the contract on EH that high-risk types like best.

Since high-risk individuals are **risk-averse** ($U'' < 0$), they prefer to equalise income across states. The marginal rate of substitution along the indifference curve equals:

$$\text{MRS} = \frac{(1 - p^H) U'(W_1)}{p^H U'(W_2)}.$$

On the 45° line where $W_1 = W_2$, we have $U'(W_1) = U'(W_2)$, so the MRS equals $(1 - p^H)/p^H$, which is exactly the slope of EH . Hence the indifference curve is **tangent** to EH at the 45° line. This tangency point α^H provides **full insurance** at actuarially fair odds — the best possible contract for high-risk types subject to zero profit.

Intuition

No one wants to pretend to be high-risk, so there is no need to distort their contract. Free to choose the best zero-profit contract on EH , risk-averse consumers prefer full insurance. This is the “**no distortion at the top**” result.

Condition on λ for equilibrium existence.

The separating equilibrium exists when λ is **sufficiently small** (i.e. there are many high-risk consumers). The reason is:

The separating equilibrium forces low-risk types to accept only *partial* insurance (the cost of separating). The alternative — pooling with high-risk types — would give low-risk types more insurance but at worse odds (the cost of pooling). A pooling deviation γ near the 45° line on the market-odds line EF attracts both types.

Whether γ is profitable depends on the market odds $\bar{p} = \lambda p^L + (1 - \lambda)p^H$:

- When λ is **small** (many H): $\bar{p}$ is close to p^H , so market odds are unfavourable. The pooling contract γ lies below the market-odds line $\Rightarrow$ it makes a loss $\Rightarrow$ no firm offers it $\Rightarrow$ **equilibrium survives**.
- When λ is **large** (many L): $\bar{p}$ is close to p^L , so market odds are favourable. The pooling contract γ lies above the market-odds line $\Rightarrow$ it makes a profit $\Rightarrow$ it breaks the separating equilibrium. And since no pooling equilibrium can exist either (Section 5), **no equilibrium exists**.

Summary

Equilibrium exists $\iff \lambda$ is sufficiently small (many high-risk). The precise threshold $\bar{\lambda}$ depends on the risk preferences $U(\cdot)$ and the gap $p^H - p^L$: the more risk-averse consumers are, or the closer together the two probabilities, the lower $\bar{\lambda}$ must be for equilibrium to exist.

Part (b): Social Planner's Pareto Improvement

What is the trade-off?

Unlike competitive firms, the social planner can offer a *menu* of two contracts (α^H, α^L) that need only break even **jointly**:

$$(1 - \lambda) \pi(p^H, \alpha^H) + \lambda \pi(p^L, \alpha^L) \geq 0.$$

This allows **cross-subsidisation**: profits earned from low-risk contracts can subsidise losses on high-risk contracts. The planner's trade-off is:

The planner's trade-off: risk-sharing vs. information elicitation

- **Better risk-sharing for low-risk types**: By using profits from α^L to subsidise α^H , the planner can give high-risk types more generous insurance. This *relaxes* the IC constraint on α^L , allowing low-risk types to buy more insurance — closer to their efficient level.
- **Cost of the subsidy**: Each low-risk individual bears a per-capita subsidy of s to fund the high-risk types' improved contract. This reduces low-risk types' consumption in every state.

The planner weighs the **gain** from moving α^L toward full insurance against the **cost** of the cross-subsidy s that each low-risk individual pays.

Concretely: suppose the planner gives each high-risk individual a per-capita subsidy of a . This shifts α^H above the high-risk fair-odds line (making losses on H). To fund this, each low-risk individual pays γa where $\gamma = (1 - \lambda)/\lambda$. This shifts α^L inward. But the improved α^H also moves U^H outward, relaxing the IC constraint and letting α^L offer more coverage. The net effect on low-risk welfare depends on whether the insurance gain outweighs the subsidy cost.

Why do high-risk consumers again get full insurance?

The planner's problem, like any optimal contract design, has the structure: choose (α^H, α^L) to maximise low-risk welfare subject to (i) IC_H : H prefers α^H to α^L , (ii) joint break-even, and (iii) H -types achieve at least some utility level.

At the optimum, α^H must provide full insurance. The argument is:

1. For any given utility level $\bar{V}$ promised to high-risk types, **full insurance is the cheapest way to deliver it**. Because H is risk-averse, any variability in their income requires a risk premium to compensate. Full insurance avoids this premium, minimising the subsidy cost.

- Formally: to deliver $V(p^H, \alpha^H) = \bar{V}$, the planner minimises $-\pi(p^H, \alpha^H)$ (the loss on the H -contract). The solution has $W_1 = W_2$ (full insurance), since the iso-profit lines of the insurer are linear while the indifference curves are convex — the tangency is at the 45° line.
- Any contract giving H the same utility with *partial* insurance would cost *more* (larger loss) and hence require a larger subsidy from L -types, which is wasteful.

Intuition

Full insurance for H minimises the cost of the cross-subsidy. Every pound of risk borne by H -types is a pound wasted (due to the risk premium), which ultimately comes out of L -types' pockets. The planner eliminates this waste by fully insuring H .

Condition on λ for a Pareto improvement to be possible.

A Pareto improvement requires that the low-risk types' gain from better insurance coverage outweighs their cost of subsidising the high-risk types. The per-capita subsidy each L -type pays is $\gamma a = \frac{1-\lambda}{\lambda} a$. This means:

- When λ is **large** (many L , few H): the subsidy cost per L -type is small (spread across many L -individuals), while the IC relaxation lets each L -type get substantially more insurance $\Rightarrow$ **Pareto improvement is possible**.
- When λ is **small** (few L , many H): each L -type bears a large subsidy burden, which outweighs the insurance gain $\Rightarrow$ **no Pareto improvement possible**; the separating equilibrium is already constrained-efficient.

Summary

A Pareto improvement on the separating equilibrium is possible when λ is **sufficiently large**. This is the intermediate range where the separating equilibrium *exists* (Section 4, λ not too large) but is *not constrained-efficient* (the cross-subsidy is affordable). When λ is very small, the separating equilibrium is already efficient and no improvement is possible. When λ is very large, the separating equilibrium doesn't even exist.

Part (c): Risk-Neutral Monopolist

Now consider a single risk-neutral monopolist (no free entry). The monopolist designs a menu (α^H, α^L) to **maximise expected profit**:

$$\max_{\alpha^H, \alpha^L} (1 - \lambda) \pi(p^H, \alpha^H) + \lambda \pi(p^L, \alpha^L)$$

subject to:

$$\begin{array}{ll} \text{IC}_H : & V(p^H, \alpha^H) \geq V(p^H, \alpha^L) & (H \text{ prefers } \alpha^H) \\ \text{IC}_L : & V(p^L, \alpha^L) \geq V(p^L, \alpha^H) & (L \text{ prefers } \alpha^L) \\ \text{IR}_H : & V(p^H, \alpha^H) \geq V(p^H, 0) & (H \text{ willing to buy}) \\ \text{IR}_L : & V(p^L, \alpha^L) \geq V(p^L, 0) & (L \text{ willing to buy}) \end{array}$$

Which constraints bind? By the standard screening logic (parallel to the monopolistic screening in the lectures):

- IR_L binds:** the monopolist extracts all surplus from low-risk types. If L had surplus left, the monopolist could raise α^L 's premium.
- IC_H binds:** high-risk types are just indifferent between α^H and α^L . If IC_H were slack, the monopolist could raise α^H 's premium further.

3. **IR_H is slack:** high-risk types earn **information rent**. Since IC_H binds and IR_L binds, we have $V(p^H, \alpha^H) = V(p^H, \alpha^L) > V(p^H, 0) = V(p^L, 0)$ (the last inequality because H values insurance from α^L more than L does from no insurance, given α^L gives L exactly reservation utility).
4. **IC_L is slack:** low-risk types strictly prefer their own contract (they would never want the overpriced high-risk policy).

The monopolist's optimal contracts:

The monopolist faces the same trade-off as the social planner but with a different objective. Distorting α^L downward (less insurance) makes α^L less attractive to H -types, which relaxes IC_H and allows the monopolist to charge H -types a higher premium (extracting more of their surplus). The marginal benefit of distorting α^L is the **information rent saved on H -types** (weighted by $1 - \lambda$); the marginal cost is the **lost profit from L -types** (weighted by λ).

When λ is low (few low-risk, many high-risk)

- Most revenue comes from H -types. The monopolist is very keen to minimise their information rent.
- The cost of distorting α^L is small (few L -types affected), while the rent savings apply to many H -types.
- $\Rightarrow \alpha^L$ **is heavily distorted** — low-risk types receive very little insurance, potentially being excluded entirely (zero coverage, $\alpha^L = 0$) in the extreme.
- α^H provides **full insurance** at a premium that extracts nearly all of H 's surplus (just above reservation utility, since IC_H and IR bind tightly when α^L is close to zero).

When λ is high (many low-risk, few high-risk)

- Most revenue comes from L -types. Distorting their contract is very costly.
- The information rent savings apply to only a few H -types.
- $\Rightarrow \alpha^L$ **is only mildly distorted** — low-risk types receive close to (but not quite) full insurance.
- α^H provides **full insurance**, and H -types earn a substantial information rent (the monopolist cannot extract it without worsening α^L , which is too costly given the many L -types).

Will high-risk consumers always receive full insurance?

Yes. Regardless of λ , high-risk types receive full insurance in the monopolist's optimum. The logic is the same as in the competitive and social planner cases:

1. The monopolist chooses α^H to maximise profit from H subject to IC_H (which determines the minimum utility H must receive).
2. For any given utility level the monopolist must deliver to H , full insurance is the **cheapest way to do it**, because risk-averse agents require a risk premium for any income variability.
3. Formally: the monopolist's iso-profit lines on H 's contract have slope $(1 - p^H)/p^H$ (the same as the high-risk fair-odds line). The tangency of H 's indifference curve with these iso-profit lines occurs on the 45° line, where $W_1 = W_2$.
4. Giving H partial insurance at the same utility level would cost the monopolist *more* (lower profit), with no offsetting benefit.

“No distortion at the top” — a universal result

In all three market structures — competition, social planner, and monopoly — the high-risk type receives full insurance. The reason is always the same: **no agent wants to mimic the**

high-risk type, so there is no incentive-compatibility reason to distort their contract away from the efficient level. And efficiency for a risk-averse agent means full insurance.

Comparison across the three regimes:

	Competition	Social Planner	Monopolist
H -type insurance	Full	Full	Full
L -type insurance	Partial	More than comp.	Partial (varies)
H -type profit	Zero	Loss (subsidised)	Positive (but rent)
L -type profit	Zero	Profit (subsidises)	Positive
H gets rent?	No	No	Yes (info rent)
L gets rent?	No	No	No (IR binds)
Cross-subsidy?	No (per-contract = 0)	Yes ($L \rightarrow H$)	Yes (implicit)
L distortion $\uparrow$ with	— (fixed)	λ small	λ small
Existence issue?	Yes (λ large)	No	No

Based on: Rothschild, M. & Stiglitz, J. (1976), “Equilibrium in Competitive Insurance Markets,” QJE 90(4), 629–649, and Moreno de Barreda lecture slides, HT 2026.