

FHS Ethics: Consequentialism

Readings

Required readings:

- Scheffler, Samuel. 1988. "Introduction." In *Consequentialism and its Critics*, 1-13. Oxford: Oxford University Press.
- Thomson, Judith Jarvis. 1984. "The trolley problem." *Yale Law Journal* 94:1395.
- Brown, Campbell. 2011. "Consequentialize this." *Ethics* 121 (4):749-771.

Optional readings:

- Sinnott-Armstrong, Walter. 2023. "Consequentialism." In *Stanford Encyclopedia of Philosophy (Winter 2023 Edition)*, edited by Edward N. Zalta and Uri Nodelman.
- Hooker, Brad. 2013. "Rule-Consequentialism." In *The Blackwell Guide to Ethical Theory*, edited by Hugh LaFollette and Ingmar Persson, 238-260. Oxford: Blackwell.
- Foot, Philippa. 1967/2002. "The Problem of Abortion and the Doctrine of the Double Effect." In *Virtues and Vices: and other essays in moral philosophy*. Oxford: Oxford University Press.

Past Year Questions

Is ex-ante or ex-post consequentialism the preferable version? [2024 I]

‘Demandingness objections to consequentialism presuppose controversial assumptions that, once revealed, undermine the power of those objections.’ Is that true? [2024 II]

‘Our reasons may depend on our antecedent ends not because those ends are the sources of the (supposed) motivating force or normative reasons, but rather because those ends are the sources of the justifying force of those reasons.’ (JULIA MARKOVITS) Discuss. [2024 II]

‘If all that killing does is prematurely end a valuable life, we might wonder why the killing especially wrongs the person who got killed, rather than just wronging all of us collectively, by removing something of value from the world.’ (CHRISTINE KORSGAARD) Do utilitarians have a good answer to this question? [2023]

Can rule-consequentialists specify a compliance level at which candidate rules are to be assessed without falling prey to decisive objections to their view? [2022]

Can consequentialists make adequate sense of moral dilemmas? [2021]

Part I: Consequentialism and Objections

Basic Overview and reasons for consequentialism

Scheffler, Samuel. 1988. "Introduction." In *Consequentialism and its Critics*, 1-13. Oxford: Oxford University Press.

- Simple Consequentialism: any right action in any given situation is the one that will produce the best overall outcome, as judged from an impersonal standpoint which gives equal weight to the interest of everyone
 - Innocent assumption of maximising the desirable and minimising the undesirable
- Two parts to a consequentialist theory
 - Theory of value: Gives some principle for ranking overall states of affairs from best to worst from an impersonal standpoint
 - Maximalist commitment: the right act in any given situation is one that will produce the highest-ranked state of affairs that the agent is in a position to produce
- Disputes in consequentialism
 - Consequentialists typically argue that our common-sense moral values which only prima facie conflict with consequentialism, can be reconciled with consequentialist principles up to some point
 - If reconciliation is impossible, then the moral principles should be revised because they are irrational
 - Non-consequentialists claim that consequentialism can at best be reconciled not with our actual values, but only with their simulacra
 - Non-consequentialists think that we need to relinquish impartiality and include agent-relative elements to reflect our actual values and that our actual values have a more secure place in human thought than abstract conception of moral rationality
- Classical Utilitarianism: the best states of affairs from among any set is one that contains the greatest net balance of aggregate human pleasure, happiness or satisfaction
 - Objection 1: utilitarianism gives no direct weight to considerations of justice or fairness in the distribution of goods (Rawls); as long as net aggregate satisfaction is maximised, utilitarianism is indifferent to how satisfactions and dissatisfactions are distributed among distinct individuals
 - Reply: Utilitarians resort to empirical assumptions that it will almost never in fact be the case that a radically inegalitarian distribution of resources or liberties will produce the greatest attainable sum of satisfaction
 - Objection 2: we are forced to do personally terrible things to do the right thing; Nagel – if the only way to prevent ten murders is to commit one

yourself, then utilitarianism requires you to do just that i.e. obligatory. In essence, utilitarianism fails to impose constraints on our actions

- Objection 3: utilitarianism is far too demanding and it requires us to neglect or abandon one's own pursuits whenever one could produce even slightly more good in some way; or it is an empty demand for in a world full of human misery and suffering only those with an extraordinary degree of moral self-confidence will be prepared to claim that there is no possible way they could do any more good for the world than by doing exactly what they are already doing
 - Williams: the demands of utilitarianism are so extreme as to undermine the integrity of the person
- Core difference between non-consequentialism and consequentialism
 - What is not a difference: non-consequentialists may agree with consequentialists in the following ways
 - There is a legitimate way of ranking the states of affairs from highest to lowest
 - Impersonal evaluations of outcomes can often affect the rightness or wrongness of actions
 - Difference are sketched in two ways
 - First, they maintain that there are occasions on which one may be forbidden to do something, even if doing that thing would have the best overall outcome.
 - I.e. the prohibited act would involved doing something horrible to someone
 - Second, they maintain that even on those occasions when one is not positively forbidden to do so, one is not always required to do it either
 - Summary: even when non-consequentialists believe that it makes sense to rank overall outcomes from best to worst, they typically believe both in agent-relative constraints (or deontological constraints) which sometimes prohibit the performance of acts that would have optimal results, and in agent-relative permissions, which sometimes makes the performance of such acts optional
 - Consists in the fact that their application to individual agents is not based or supported by any appeal to what would be best from an impersonal standpoint
- Consequentialist rejoinder: not necessary to accept a version of agent-relative morality in order to avoid the difficulties encountered by simple utilitarianism
 - To a large extent, utilitarianism's vulnerability on the issue of distributive justice can be attributed to the specific way in which it evaluates outcomes
 - Given any two outcomes with different totals of aggregate satisfaction, utilitarianism will always say that the outcome with the higher total is better

- Many consequentialists would agree that this is an implausible way of evaluating outcomes → natural solution is to substitute a principle of evaluation that is more sensitive to distributive concerns, rather than abandoning consequentialism altogether
 - Two ways to accommodate distributive concerns (Scanlon)
 - Giving extra weight to the interests of those who are worst off, so that the satisfaction of their interests count disproportionately in determining what the best outcome would be in any given situation
 - By treating distributive equality as a good in itself, which must be considered along with factors like net aggregate satisfaction in determining the value of an overall outcome
 - Difficulty is with the integrity objection because it is fundamental to achieving the best outcome
 - Railton: Alienation, Consequentialism and the Demands of Morality → the right thing to do in these cases is, as always, to perform the act that will produce the best overall outcome
 - Rule utilitarians
- The second objection is also potent: there are occasions when one must not violate an agent-relative constraint even if that is the only way to prevent more widespread violation of the very same constraint by others
 - Nozick's side constraints
- Deontological Constraints: the task
 - Victim-based explanation: The explanation can proceed by identifying something about people that makes it inappropriate to victimise them in the ways that the constraints prohibit
 - But this tends to fail to explain why we may not victimise one person who has that property in order to prevent the victimisation of an even larger number of people, each of those who has that very same property
 - Agent-based explanation: point to something about the agent that makes it wrong for him to victimise one person even in order to minimise the total overall number of comparable victimisations
 - Overall: strong doubts about the rationality of agent-relative morality continue to be expressed
 - But we are also unwilling to abandon our deeply felt agent-relative convictions, or to submit to them to consequentialist reinterpretation
- Overall impasse: On one hand, there are some insisting that agent relativity is irrational and others insisting that our agent relative intuitions are so strong and entrenched that it would be a mistake to abandon them
 - Potential Solution 1: takes as the starting point the observation that the sharpest apparent conflicts between agent-relative morality and our

conception of rationality are generated by agent-relative constraints – by the claim that one is morally forbidden to minimise morally objectionable activity

- Suggestion is: we should examine the possibility of a moral view that departs from consequentialism to the extent of incorporating agent-relative permissions, but not to the extent of accepting agent-relative constraints
- On a view of this kind, one would always be permitted but not always required to perform the act that brings about the best available outcome
- Potential Solution 2: Sen argues that (i) we have to incorporate rights as features of states of affairs which must be taken into consequential evaluation, (ii) we have to make certain finer-grained distinctions among different types of agent-relativity, (iii) we must allow rankings of overall outcomes to vary from person to person
 - Possible range of moral views widens considerably
 - Foot takes a similar path, by arguing that consequentialism rests on the assumption that it always makes sense to talk about the goodness or badness of the state of affairs produced by an action (without this, consequentialists could not formulate the charge that agent-relative moralities sometimes tell us, irrationally, to choose a worse state of affairs over a better)
 - Foot argues that the assumption is false: expressions like ‘best state of affairs’ makes sense only in certain limited contexts where the virtue of benevolence provides them with a sense
 - Scheffler rejects that consequentialist depends on this assumption → the difficulty is deeper
- Potential Solution 3: sophisticated forms of consequentialism to try to unify non-consequentialist and consequentialist views
 - Johnson argues that Rawls and Raz to arrive at a more fully satisfying reconciliation

Alienation Objection

What the problem is

- Utilitarianism forces us to make decisions in a way that seems contrary to being human and so is intolerably negatively self-affecting
 - The highly rational, utilitarian way in thinking about what we ought to do alienates us from our special commitments because it abstracts us away into an **impartial spectator** when we decide what to do
 - It neglects that we have agent-relative reasons for pursuing our personal prerogatives and special relationships

- The alienation objection is therefore the **consequence** of the explanatory problem in utilitarianism
- Decision procedure problem: Self-alienation aspect
 - Railton: By employing universal principles in our decision procedures and our moral deliberations, we regard others in ways that alienate them from our considerations and relationships
 - Counterintuitively, we seek justification for our actions to help them not by virtue of them or our relationships, but from the cold, abstract utilitarian principle
 - Important form of alienation in moral practice (applied in both decision-procedure and criterion): the sense that morality confronts us as an alien set of demands, distant and disconnected from our actual concerns
 - **John and Anne:** John is the model husband to Anne and shows great sensitivity towards the needs of Anne. He feels great affection for her. But he justifies his actions by saying “I have always thought that people should help each other when they’re in a specially good position to do so. I know Anne better than anyone else does, so I know better what she wants and needs. Besides, I have such affection for her and it is no great burden. I get great satisfaction out of it. Just think how awful marriage would be, or life itself, if people didn’t take special care of the ones they love.”
 - John demonstrates alienation: there is an estrangement between their affections and their rational deliberative selves
 - **An abstract and universalising point of view mediates their response to others and to their own sentiments**
 - Continuous with Mill’s application of a general rule to a particular instance: nothing inherent about that particular instance and thus cannot accommodate special relationships
 - John is not necessarily suffering in the utilitarian sense from this lack of alienation; it simply contradicts how we conceive of humans as agents **already situated** within social circles and bounded by personal commitments
 - So, the way in which utilitarianism traditionally forces us to make decisions alienates us from our personal relationships and special commitments, **inhibiting us from being able to deeply engage in these relationships in a way that makes us distinctly human**
- Criterion problem: Explanatory inadequacy
 - Source of disagreement: AUs think that we should be a cool and rational calculator of welfare, with no personal attachments, affections or concerns and so evaluates rightness based on such an emotionless model of humanity

- Williams offers a contrary, emotionally engaged agent who allows themselves to develop deep partial attachments to others, and who lives by or at least consults emotions in moral situations
 - Such a person may create special constraints against the unthinkable (killing) or may pursue person relationships for emotional and sentimental reasons
 - They allow their emotions and sentiments to influence their actions directly
- Crisp argues that at the level of theory, it is not as if an act utilitarian must treat utilitarianism as completely independent of emotion
 - Rather, Crisp finds that it is open to the utilitarian to argue that our sentimentality and moral emotions towards our loved ones are **justified by utilitarianism**
 - Mill's proof is this kind of argument, by saying that the utility principle is what tacitly influences our common sense morality and our intuitions
 - However, this is **merely a theoretical possibility**; it seems more likely that utilitarianism does not permit such coherence with our sentimentality
- The more important point is that utilitarianism fails to give sufficient weight to moral emotions and the reasons for why we pursue certain actions that they reveal
 - It seems counterintuitive that in cases such as loving your wife or caring for your children, the reason for your action is utility maximisation, rather than embodying the virtue of love and caring for your children because for their sake
 - The inherent explanatory inadequacy within utilitarianism is not only that it cannot square off against certain intuitions, but that it provides **reasons for actions that are alien** to the moral agent in the case of special commitments
 - Utilitarianism confronts us with an alien set of demands, distant and disconnected from our actual concerns
- Note that alienation is not always a bad thing
 - Often desirable that people achieve some distance from their sentiments of one another

Solution

Railton's sophisticated consequentialism (essentially MLAU)

- Basically, Railton proposes that we dispense with the act utilitarian decision procedure and instead look towards what kinds of people are the happiest
 - Note that they have strong loyalties to friends

- So dispensing the AU decision procedure is **consistent** with objective hedonism as a criterion of rightness
- We **do not need to be committed to consequentialist deliberations** and the agent can fully recognise that he develops his dispositions because they are necessary for promoting the good
- Objective consequentialism has the virtue of not blurring the **distinction** between the **truth-conditions** of an ethical theory and **its acceptance conditions in particular contexts**
- I depart from Railton in thinking that subjective hedonism must be dismissed
 - In fact, in observing what kinds of people are happiest, I am not making the claim that if I were to do what they do, I would definitely be happiest; rather, I am making the claim that if I follow their deep engagement in their commitments, I would probably be very happy, such that I maximise expected happiness
- Railton's counterfactual condition: the sophisticated hedonist may not always act for the sake of happiness, since he may do various things for their own sake or for the sake of others. But he would not act as he does if it were not compatible with leading an (objectively) hedonistic life/it does not maximise utility
 - So if a devil told someone that he had to be fully Kantian, otherwise he would destroy the world, the agent adopting a fully Kantian perspective would still meet the counterfactual condition and qualify as a sophisticated hedonist because he would not have adopted Kantianism should it not have been the utility maximising option
 - How far removed can the criterion of rightness be from our chosen decision procedure and motivational structure?

Adopting a plurality of ends beyond hedonism (Railton)

- A pluralistic approach in which several goods are viewed as intrinsically, non-morally valuable — such as happiness, knowledge, purposeful activity, autonomy, solidarity, respect and beauty
 - One mistake of dominant consequentialist theories is their failure to see that things other than subjective states can have intrinsic value → desire theory & hedonism
 - Related is the misconception that we need to reduce all intrinsic values to one — happiness
- First, in divorcing subjective states (the happiness produced by the action) from their objective counterparts (the actual action), and claiming that we seek the latter exclusively for the sake of the former, utilitarianism cuts us off from the world in a way made graphic by the experience machine

- The experience machine affords us decisive subjective advantages over actual life: few, if any, in actual life think they have achieved all they could want → but most rebel against it
 - It matters what we actually do and are as well as how life appears to us
- Second, the reduction of all goals to the purely abstract goal of happiness and pleasure treats all other goals instrumentally
 - Knowledge and friendship may promote happiness but is this a fair characterisation to say that this is the only way in which they are valuable?

Problems with Railton's solution

- Does not address the explanatory objection
 - As per usual
- Open to empirical objection; it does not guarantee that this sophisticated variation of consequentialism will adopt such a non-alienating decision procedure
 - Railton: Possible that the happiest sort of lives ordinarily attainable are those led by people who would reject even sophisticated hedonism, people whose character is such that if they were presented with a choice between two entire lives, one of which contains less total happiness but nonetheless realises some other values more fully, they may choose against maximal happiness
 - It may be the case that RU is the optimific set of rules upon internalising the cost of the motivational structure → a difficult, open empirical question
 - In fact, Railton initially points out that the alienation problem may not cause a reduction in consequentialist value at all → even less likely for a sophisticated hedonist's motivational structure to be adopted
- Standard issues with OLT: aggregation problems, explanatory problems and ironically, alienation problems
 - How can we claim that OLT reduces alienation in our moral theory, when it plausibly confronts us with an **alien definition of what is good for us**

Integrity/Demandingness Objection

Version 1: No agent-relativity

- Bernard Williams: Utilitarianism leaves no room for us to accommodate our special projects, person pursuits and prerogatives, and thus cannot accommodate our integrity
 - Crisp coins an analogy that Williams means integrity as **integrity in a work of art**: Artistic integrity is generally defined as the ability to omit an acceptable level of opposing, disrupting, and corrupting values that would otherwise alter an artist's or entities' original vision in a manner that violates their own preconceived aesthetic standards and personal values

- So basically, our capacity to preserve our own sense of self; our vision of who we want to be and how our life should be
 - This is constituted by our deeply held moral principles, our social and relational commitments and our personal projects → this defines the individuality of our lives
- Aspects of integrity: essentially having to do with **preserving the sphere of agential authority**
 - **Principles:** Forces us to violate our deeply-held principles
 - Jim and the Indian and George and the chemical weapons
 - **Personal prerogatives:** Forces us to give up our personal prerogatives and often deems the pursuit of personal relationships impermissible
- Aspects of the objection (decomposing the objection)
 - **Substantive objection:** does it not matter morally in certain situations that it is you in particular who are acting in a certain way; are there not constraints on what you can do
 - **Explanatory objection:** utilitarianism requires us (through its decision procedure and moral thinking) to employ an unrealistic notion of a self that can float free of specific commitments and concern itself solely with welfare maximisation
 - So it begins from an unrealistic (and thus false) standpoint of the unencumbered self; if it begins from a flawed premise/account of human nature, it follows that the explanation of morality is flawed as well
- Key intuitions that we must consider
 - TE1: George being offered a job to create chemical weapons. He is morally objected to this but is told if he does not take up the job, someone enthusiastic would take his place and thus lead to far more developments of chemical weapons that he would allow.
 - Williams suggests that act utilitarians would advance that it is *obvious* that George should take the job
 - However, it does not seem obvious that George should take the job; intuitively, we would not blame George for not
 - TE2: Jim being forced to shoot 1 Indian so that 19 live by an evil captain vs him not doing anything and the captain shooting all 20
 - Again, act utilitarians deem it obvious that Jim shoot the single person
 - But given that Jim has deeply held principles against killing people, utilitarianism's demands would force him to violate his integrity
 - Intuitively therefore, we think that if Jim does not kill the single person, he has done nothing wrong since the preservation of his integrity is his personal prerogative (it is within the scope of his agential authority)

- Lessons from the intuitions
 - Do we accept the utilitarian conclusion? Substantive considerations
 - Utilitarianism does not seem to allow for agential authority/personal prerogatives and neither does it seem to allow for special constraints in what we ought to do
 - Do we accept the way that utilitarians look at these questions? Explanatory considerations
 - Perhaps Jim should kill the single person but not because it is of maximal utility, but because it aligns with some virtue of being able to make sacrifices/bravery
- Substantive objections
 - Ignoring special constraints to what we can do
 - Excluding the existence of personal prerogatives, agential authority and agent-relative reasons in pursuing certain ends
- Explanatory objection
 - Ignores agent-relative reasons for pursuing certain ends
 - Williams: utilitarianism can only make the **most superficial sense of human desire and action**
 - We ignore our own characters, desires, goals, projects and commitments and ask not 'What should I do?' but 'What does utilitarianism require of any person so situated in this position in order that the history of the world go as well as possible'?
 - Our moral responsibility is a result of coincidence → I am just a representation of the satisfaction system who happens to be near certain causal levels at a certain time → does not matter that the moral agent is you specifically
 - So utilitarianism is wrong in its explanation of rightness as maximising the good, because this **undermines our individuality** inextricably tied to being human
 - Crisp: utilitarianism provides **too narrow** a view of the reasons for human action; by supposing that maximising utility is the only reason for moral action, it neglects agent-relative reasons such as special relationships
 - Imagine Smith who has a friend in the hospital but has a crucial work meeting that would provide more utility overall
 - Even if utilitarianism is right in thinking that going for the work meeting is the right action, we should still think it plausible that Smith has a **special reason to visit you just because you are her friend** (existence claim)
 - Godwin: there is indeed a magic in the pronoun "my"

- And if visiting her friend is the right action, it seems to follow that maximising utility is not the right action, but her special commitment to being her friend
- Utility maximisation seems to **be the wrong reason** for pursuing our ends; there seems to be cases that we ought to do the action with suboptimal utility or even negative value
 - Imagine Smith is the president of an NGO and is running a considerable gala dinner to raise funds. Donors are excited to meet her and she would be able to garner moderately more funds for the charity. She just received notice that her close friend has gotten into a car accident and is comatose. She is going to die. There is 0 value in Smith visiting her. But does she still not have a special obligation to do so, over the dinner?
 - **The Father.** A father's son has leukaemia and is about to die in a week. He is comatose. The father plays a crucial role in an innovative AI company that betters the world. He knows that if he visits his son this last week, it will cause him so much anguish and pain that he will not be able to go back to work for a month. He will likely get fired and the company's next innovation will be significantly delayed.
 - Utilitarianism would suggest that the right thing for the father to do would be to not visit his son
 - Runs counter to two intuitions
 - The first is that regardless of what the father does and how much value we ascribe to either action, we do not think that he has done anything wrong. The father's special commitment is his personal prerogative – he ought to have autonomy over whether to fulfil his commitments.
 - The second observation is even stronger: Some may say that realising fatherhood has significant value, and thus may be why the father visiting his son may be right. But it seems that the lower the scenario's capacity for value promotion, the more significant fulfilling the special relationship is. The fact that the son is comatose, that the father is causing more pain and anguish, engenders a greater significance of fulfilling that special relationship. Even if realising fatherhood has value, the decrease of net welfare, on utilitarian grounds, should not incline us to morally defend the father even more. So, it cannot be right that we deem his actions *immoral*
 - Two conclusions from this
 - First, the importance of our special commitments is not based on how much consequentialist value they have. Utilitarianism

makes an explanatory mistake in *why* we consider our special commitments important. It follows that utilitarianism cannot correctly and possibly justify our special commitments as right since they are important respective of consequentialist value; rather, our special relationships seem to possess some non-consequentialist value (bound to beg the question against consequentialism)

- So utilitarianism's wholly consequentialist foundation seems untenable, because it excludes important non-consequentialist considerations and value embodied by our special commitments and personal prerogatives
- Second, what matters is not whether the father is acting in the right way. What matters is that the father ought to have agential authority to fulfil his personal commitments, and this needs to be considered permissible by utilitarianism

Solution

Jackson's expected utility solution: Focuses on examining how utilitarianism's criterion of rightness can account for why we think pursuing our personal prerogatives and special commitments are morally right

- To Jackson the criterion of rightness must adopt a probabilistic notion of value, rather than an actualist notion of value
 - What this essentially means is that the right action is decided *ex ante facto* (prior to the actual consequences) based on the probabilities of each outcome for that particular action, rather than *ex post facto*, where the action is evaluated based on the actual consequence that arises
 - So the right action is determined by which action maximises the expected value where the expected value of each action j is determined by the summation of objective probability function $Pr(O_i|A_j)$ multiplied by the consequentialist value derived from that outcome for each outcome i , given action j

Rule of action:

$$\max \sum_i Pr(O_i|A_j) \times V(O_i)$$

$Pr(O_i|A_j)$:

agent's objective probability function of outcome

$V(O_i)$:

objective value function of outcome i

- The nature of the probability function is **agent-neutral** in the sense that the expected value of each action's outcome is **not affected by agent-relative** reasons such as greater care for his special commitments; it is **only affected by the information** the agent has regarding the likelihood of each outcome
 - Note that since the probability function is contingent on the agent's cognitive ability and information the agent has, it is necessarily subjective in the sense of **being restricted to the moral agent's viewpoint**
 - Despite being subjective, it remains agent-neutral in a way that is consistent with utilitarianism's requirement of **impartiality**: that everyone's utility counts for the same
 - This impartiality is captured by the objective value function $V(O_i)$ which simply takes any outcome i as its argument, and churns out some utilitarian value **based solely on its consequentialist calculation methodology**
 - So the value function $V(O_i)$ is identical to an actualist utilitarian value function – Jackson merely multiplies it by a probability function to transform actual value into expected value
 - The consequentialist value function assigns no value to the intention of maximising expected utility
 - Our decision theory has no intrinsic value on the consequentialist account → its only value is its tangible welfare impacts
 - The subjective probability function is an idealisation of the agent's beliefs in the sense that they are not misinformed but accurate
 - There exists the possibility of culpable ignorance
- Reasons why the probabilistic criterion of rightness is adopted
 - A1: The criterion of rightness should be limited by an agent's epistemic constraints (intuitionistic argument)
 - An agent cannot possibly know which of the probabilistic outcomes
 - Consider the case of a doctor with the option between a drug that has a 100% chance of partially curing a patient vs a drug that has a 90% chance of killing a patient and a 10% chance to completely cure the patient
 - Even if the doctor chose the second drug and it turned out to have been successful, intuitively, we would think that the doctor chose the wrong action because he was more likely to cause significant damage than not. Thus, the doctor, in not maximising expected value, had **no proper reason to choose** the riskier drug other than pure, baseless hope

- Adjacently and less significantly, it circumvents the problem of moral luck (which explains our intuitions)
 - When we give moral praise, we usually do so because people are responsible for the good deed
 - In the doctor case however, given that probabilistic outcomes are exogenous to our actions, he was not responsible for the good deed since it was completely unexpected that it gave a superior outcome
 - In effect, he did not commit the good deed, “probability” did
- A2: If the criterion of rightness is independent of an agent’s epistemic constraints, then the moral theory is self-defeating by not giving a proper guide to action
 - Suppose consequentialism says nothing about the mind of the agent: merely that the right action is action with property X, for some consequentialist treatment of X which pertains solely to what in fact would happen and not at all to what the agent thinks
 - Williams makes the point that consequentialism must say something about the right decision → objective consequentialism gives no figment of what the right guide to action is (though it creates a standard for us to test our guides against) since what actually happens is subject to probability
 - On the other hand, if consequentialism is expressed in its subjective variation and about how we are supposed to make moral decisions and it says to think along X lines. What if then thinking along those lines is discovered to have bad consequences in certain situations
 - Jackson proposes that his decision theoretic account of consequentialism disarms the second horn of the dilemma by answering that in such situations, the agent ought not to think along those lines (because after all he commits to subjective probability) → why does objective consequentialism not allow for this non-committal
- How this can solve the integrity/demandingness objection
 - Thus conceived, Utilitarianism now allows us to help those dearest to us (our loved ones) as opposed to those distant from us because the probability of success in helping them (given that we have more knowledge regarding their needs and wants) is significantly higher
 - So, it may be the case that the expected value is maximised by pursuing our special commitments and person prerogatives
- Problems with the solution
 - Incompatible with modern developments
 - Singer points out that with modernisation, information technology and instant communication, our epistemic constraints are not as

severe as Jackson assumes – we can know the needs of the children starving in Africa by looking it up

- Even if one counters by pointing out that interpersonal connection provides more intimate knowledge of the situation than internet based knowledge, it is obvious that the difference between how both types of information elevates your capabilities in helping each group of people is miniscule
- Moreover, it tends to be the case that for affluent people, the raw value that you can produce by helping those distant from you (suffering children) is infinitely greater than the raw value you can produce by helping those nearest to you
 - It tends to be that those nearest to you face troubles that seem undeniably trivial compared to the starvation that the poor face – they teether on the edge of death
 - And if we accept that helping one person live for a few more months produces much much greater value than helping your neighbour when his car breaks down, it seems difficult to consistently accept that when adjusting for expected value, it turns out that helping your neighbour produces more expected value than helping the poor
- Narrow in scope; neglects significant aspects of the integrity objection
 - What about the example of Jim and the Indian when one's deeply and reasonably held principles are being violated by utilitarianism?
 - Jackson's solution only deals with a specific class of the integrity objection, relating to special commitments and personal relationships
 - It may not even be extended to salvaging the pursuit of personal projects because the value behind these acts of self-realisation are often minimal compared to what value you could provide to others distant to you
- Cannot overcome the explanatory objection
 - Does not address the concern that utilitarian reasoning seems to be intuitively the wrong reason for pursuing personal commitments and relationships → in fact, it seems to exacerbate it by perpetuating the cold, utilitarian rationality in justifying our pursuit of personal commitments
 - See father example in trying to prove this point
 - Even if one adopts some sophisticated utilitarianism and does not act on the basis of utilitarianism when engaging with his special commitments, we seem to be committed to the fact that that specific decision procedure the agent used (that he loves his wife) is the

correct reason for celebrating her birthday and is **why the action is right**, and not because of some utilitarian/consequentialist reasoning

Denying that integrity is an actual objection because it is self-contradictory

- On one hand, integrity emphasises the need for a sphere of agential autonomy because of our personal projects and special commitments
 - But autonomy mandates that we need to be able to critically evaluate the value of these projects and commitments → if we shield them from evaluation, we are essentially claiming that, to an extent, are personal prerogatives are morally infallible → obviously not the case
- So it may not make sense for us to speak of “integrity” because our projects and commitments should always be up for scrutiny → makes no sense for us to preserve our commitments for the sake of it; morality grants us a framework to evaluate them
 - But of course, morality should not be so destructive as to remove ourselves from everything
 - Consider also that if morality is universal and forces us to reevaluate our projects, would our new projects not be convergent with one another? So actually, morality disables our autonomy by dictating that our projects ought to converge with one another
- Should not exaggerate autonomy to be the Kantian conception of escaping determination by any contingency or whatsoever + refuse to conflate autonomy with sheer independence from others

Version 2: Negative Responsibility

- Doctrine of negative responsibility: ignores the distinction between **causing** something to happen and **allowing** something to happen
 - The problem is not that utilitarianism thinks that allowing people to die and causing them to die are both wrong, but rather that utilitarianism is indifferent between both scenarios, ascribing **equal** moral (dis)value to both actions
 - After all, Peter Singer argues that if we walked past a pond and see a child drowning, we ought to help him should it pose no significant cost to ourselves
 - The problem is that utilitarianism does not appreciate that causing someone to die is worse than letting him die
- Another aspect of negative responsibility is that utilitarianism becomes a rather unfair moral doctrine because it makes you **responsible for what other people do**
 - First, it punishes seemingly innocent agents for other’s immoral actions
 - Consider the fact that the world would be a better place if everyone donated 1% of their income to the poor

- As more and more people violate this moral rule, those who have donated would be required to donate even more of their income (say 10% instead of 1%)
 - It seems completely unfair that because other people are immoral and do not donate their income, I have to now donate more of my income despite having donated my corresponding amount
 - Singer objects to this intuition: consider the scenario that there are 2 boys drowning and both you and your partner ought to now save one boy each. Imagine your partner does not save his boy. You now ought to save both boys. Even if we accept that this seems unfair to you, one could not deny that saving both boys is the right action for you to undertake
 - It could be that morality by nature is demanding and unfair
- Rest is the same as integrity because in forcing us to violate our integrity, utilitarianism is demanding

Solution

Sobel's impotence of the demandingness objection

- Basically, a non-demanding theory on moral agents/benefactors is demanding on the sufferers.
 - He further argues that we must look towards a distinction between the moral value of a theory **causing compliance costs** on its actors vs **allowing costs to happen** on the recipients of actions.
 - This is similar to whether there exists a moral distinction between requiring that costs happen on the benefactors vs letting these costs happen on the beneficiaries
- Consequentialism implies that only outcome values matter, and thus adopts the doctrine of negative responsibility: consequentialism is indifferent between whether someone let these outcome values occur or caused these outcome values to arise + is indifferent between who is responsible for these outcome values → the action or responsibility of the agent is secondary to the occurrence of the outcome value
 - Thus, the demandingness objection requires us to break from these consequentialist assumptions before launching the attack. We have to deny consequentialism before constructing the demandingness objection, implying that the demandingness objection itself is not a fundamental objection.
 - The demandingness objection is in effect begging the question against consequentialism and utilitarianism
- The main justification for the demandingness objection is not simply that the compliance costs of utilitarianism are too high, but that these costs are those that we are morally entitled not to pay

- This means that it supersedes any utilitarian or consequentialist calculation
- To say that there is a significant moral distinction between requiring and permitting is just to say that a moral theory is less accountable for what it permits than what it requires
 - But consequentialism holds one fully accountable for what one allows will be too demanding with respect to what one may allow
 - So one must object to the doctrine of negative responsibility

Denying the uniqueness of this charge (particular for negative responsibility)

- Negative responsibility as a doctrine may not be exotic
 - Singer points out we ought to save the person drowning in the pond
 - Many legal systems adopt such a doctrine in punishing people who knowingly allow something bad to happen
- The key problem is that utilitarianism treats causing something bad to happen and allowing something bad to happen with complete indifference
 - Surely, intentionally killing someone is worse than letting someone die?
 - On retrospect, this doctrine seems largely acceptable
 - Killing someone seems marginally worse than letting someone die → is this simply a result of our psychological dispositions and not actually a reliable intuition

Objections from Rights

Thomson, Judith Jarvis. 1984. "The trolley problem." *Yale Law Journal* 94:1395.

Central thesis: when facing competing intuitions over the trolley and hospital problem, Thomson suggests a rights-based principle to differentiate the two i.e. we have to invoke rights to make sense of our intuitions

- Central Problem
 - Trolley: if a trolley comes hurtling down towards five track workmen who have been repairing the track, do you as the trolley driver turn the trolley that the trolley hits a single track workman
 - Most people say that it is morally permissible to switch the track
 - Some say it's morally required
 - Hospital: Imagine you are a great surgeon who is 100% successful in transplant surgeries and that there are 5 people dying. Coincidentally, you can harvest all the organs needed from one man, and there is only one man available and he is alive. Are you allowed to kill that one man to save 5?
 - Most people forbid such an action
 - On a consequentialist/utilitarian thesis, there is nothing morally different between the two cases
 - Foot: Why is it that the trolley driver may turn his trolley while the surgeon may not remove the young man's lungs, kidneys and heart

- Foot's solution (Killing vs letting die distinction): The surgeon and the trolley driver face different decisions
 - Killing one is worse than letting five die
 - It is the surgeon's choice between killing one (operating) and letting five die (not operating)
 - Killing five is worse than killing one
 - The trolley driver's choice is between turning the trolley and killing either group of people
 - Thomson accepts that for these cases as described, the killing vs letting die distinction holds but argues that it does not generalise
- Foot's solution does not generalise
 - Bystander at the Switch: instead of being the trolley driver, you see the switch and can switch the trolley's direction
 - You now have a choice between throwing the switch (killing one) or not throwing the switch (letting five die)
 - Thomson thinks that the intuition is we still think it is morally permissible to throw the switch
 - What exactly is the difference between the driver failing to turn the trolley and the bystander failing to throw the switch when the driver would just be sitting there and doing nothing
 - Here the killing and letting die distinction seems contentious
 - Thomson just advocates for the weak thesis that the bystander may intervene → and if this is true, it shows that Foot's solution does not generalise
 - So killing one is worse than letting five die is false generally
- The Killing five is worse than killing one doctrine may also meet problems
 - Careless surgeon: Imagine that the surgeon is careless and administers a fatal drug into 5 patients, that damages each patient differently via different organs, who now require transplants. The surgeon now has a choice of cutting up one man to save the 5 (and thereby killing one instead of 5)
 - But perhaps there is something to do with the cause of death
 - Does not seem permissible for him to cut up the man
 - Some killings themselves include lettings die → where the act by which the agent kills takes time to cause death – time in which the agent can intervene but does not
- Thomson thinks that the doctrine of "killing five is worse than killing one" is true, *ceteris paribus*, but it also does not tell us that the surgeon ought to operate
 - May shift the attention to "If a person is faced with a choice between doing something here and now to five, by the doing of which he will kill them, and doing something else here and now to one, by the doing of which he will kill

only one, then *ceteris paribus* he ought to choose the second alternative rather than the first

- Trolley driver is confronted with the here and now while surgeon is not
 - Surgeon merely fails to do something by the doing of which he would make it the case that he has not killed five
- But why does the present tense matter so much: why should a person prefer killing one to killing five if the alternatives are wholly in front of him, but not where one of them is partly behind him?
- Kantian approach is unsatisfactory: Act so that you treat humanity, whether in your own person or in that of another, always as an end and never as a means only
 - Surgeon ostensibly would treat the man as a means only while Bystander at the Switch does not use his victim to save five
 - But not clear what it consists of to “use as a means only”: in some counterfactual, if the other track disappeared, then the bystander would be forced to let the trolley hit the 5
- Two facts which explain what is morally different between the agent in the Hospital/Transplant case and the Bystander case
 - The Bystander saves his five by making something that threatens them instead threaten one
 - In the surgeon case, it is organ failure that threatens the five, while it is cutting the one up that threatens him
 - Would be permissible to redirect poisonous gas to that one person
 - The Bystander does not do that by means which themselves constitute an infringement of any right of the one’s
- Concept of a right: Dworkin say that “rights trump utilities”
 - Maximising utility is not a sufficient reason to override one’s rights → Nozickian side constraints
 - The surgeon would be infringing on the man’s rights and the young man has a right to life
 - The bystander may proceed because he maximises the utility and in so doing does not infringe any right of the one track workman’s
 - But it is not clear that the bystander would infringe no right of the one track workman’s if he turned the trolley
 - But there is some reason to think that the bystander will infringe a right of the one if he throws the switch → so the explanation from rights may not be sufficient if we seek to hold that the bystander may proceed
- Thomson’s solution
 - Why is it morally important that the same thing threatens both in the Trolley case whereas it is not so in the Hospital case? *Ceteris paribus*, to kill a man is

to infringe his right to life and we are prohibited from killing; it is not enough to justify killing a person that if we do so, five others will be saved

- But that five others will be saved is not all that can be said in defense of killing in Bystander – the bystander who proceeds does not merely minimise the number of deaths which get caused – he minimises the number of deaths which get caused by something that already threatens people, and that will cause deaths whatever the bystander does → but why does the same thing causing death matter?
- Distributive exemption: the bystander does not make something be a threat to people which would otherwise not be a threat to anyone
 - Permits arranging that something which will do harm anyway shall be better distributed than it otherwise would be
 - Not just any distributive intervention is permissible: it is not generally morally open to us to make one die to save five, but *ceteris paribus*, it is not morally required of us that we let the burden descend out of the blue onto five when we can make it instead descend onto one
- Principle: it is not morally required of us that we let a burden descend out of blue onto five when we can make it descend onto one if we can make it descend onto one by means which do not themselves constitute infringements of rights of the one
 - Fat man: you can push the fat man to stop the trolley from killing the five → but you can't do so because he has a particularly stringent right against it (Thomson thinks this holds even if he does not die)
 - Objection: suppose you do not need to push the fat man but you can wobble the handrail to drop him onto the track – this seems analogous to the Bystander case → the means you would be taking to make the trolley threaten one instead of five does not constitute an infringement of a right
 - Reply: tighter notion of “means” → by hypothesis, wobbling the handrail will cause the fat man to topple onto the track in the path of the trolley, and thus will cause the trolley to threaten him instead of the five → but the trolley will not threaten him instead of the five unless wobbling the handrail does cause him to topple
 - Getting the trolley to threaten the fat man instead of the five requires getting him onto its path
 - The means you take to get the trolley to threaten the one instead of the five include wobbling the handrail, and all those further things that you have to succeed in doing by wobbling the handrail if the trolley is to threaten the one instead of the five
- Not just a right of the fat man, but a particularly stringent right of his

- Mayor case: suppose that the Mayor invites people to have a picnic on the right-hand track, which has been unused for 10 years, promising that there will not be a trolley. Suddenly, a trolley comes throttling down towards the left-hand side track and the Mayor has a choice
 - Controversial
- Thomson admits that there will be cases in which it is vague whether a means constitutes infringement of stringent rights
- Assessments of which acts are worse than which have to be directly relevant to the agent's circumstances if they are to have a bearing on what he may do
 - If A threatens to kill five unless B kills one, then although killing five is worse than killing one, the alternative is not that B kills five or B kills one but rather A kills five or B kills one → judgement that it is worse for B to kill the one is compatible with killing five is worse than killing one

Objection from liberty rights

Summary: Both Rawls and Nozick's objections are fundamentally explanatory in nature: they try to explain why utilitarianism **cannot capture our substantive notions and convictions** about justice by arguing that utilitarianism relies on **flawed premises**

- Rawls' separateness of persons critique
 - Utilitarianism's teleological foundations leads it to define the right in terms of the good, and define the good independently and prior from the right
 - So the right cannot set constraints which are independent from the good
 - In contrast, justice as fairness is deontological which defines the right independently from the good → the concept of the right is prior to that of the good; in this sense, the right is designed to set constraints on the pursuit of the good: "impose restrictions on what are reasonable conceptions of one's good"
 - Utilitarianism undermines the separateness of persons by extending the rational principle of one man to freely distribute and maximise utility along different points in his life to that of society; the utilitarianism therefore thinks that society can similarly unboundedly reallocate utility across persons
 - Utilitarianism wrongly conflates different people into a single man
 - This mistake arises from the fact that while the man holds a single life plan at any point in time, society is constituted by discrete individuals with different life plans at any one point in time
 - It does not seem just that one person ought to sacrifice the realisation of his life plan for another's who is wholly unrelated to him
 - Society is constituted by a plurality of distinct persons with separate systems of ends; it follows that each person's system

of ends are detached from one another in that one person's loss **cannot be compensated** by another's gain

- Our individuality and individual life plans ought to be protected such that there is no reason to think that my loss can be compensated by another's gain; for in order to justify such transfers, one must have a **right to my compensation to him, and I must have a duty to compensate another**
 - The utilitarian reason that "he has a right to your compensation because it is utility maximising" seems **insufficiently strong to constitute a right** because I am not actually compensating him for anything
 - If we think that such claim-rights can be motivated by such weak reasoning, then the whole conception of personal ownership will be lost
- So a person's liberty right such as right to property and right to life can be easily sacrificed for the sake of greater overall utility
- Nozick's side constraints critique
 - Nozick's critique closely follows Rawls. He objects to the idea of society employing the rational principle of one man to redistribute utility amongst people. In fact, he goes much further in arguing that there actually does not exist a social entity
 - But there is no *social entity* with a good that undergoes some sacrifice for its own good. There are only individual people, different individual people, with their own individual lives. Using one of these people for the benefit of others, uses him and benefits the others. Nothing more. What happens is that something is done to him for the sake of others. Talk of an overall social good covers this up
 - Nozick argues that what makes it rational for someone to incur a cost now is to avoid incurring a greater cost
 - But for society, it is someone incurring the cost without the avoidance of a greater cost or some overall net gain in his personal good
 - Nozick begins from Kant's premise that people ought to be treated as ends in themselves
 - Such unbounded redistribution of the good amongst people treats people as means to some societal (if such a thing even exists) end or even to the ends of others
 - Nozick maintains that the endorsement of the imposition of sacrifices on individuals for the sake of *any* (non-trivial) conception of the social good—even, e.g., distribution-sensitive conceptions of the social

good—fails to recognize the status of individuals as ends-in-themselves.

- Nozick makes the **substantive point** that what utilitarianism gets wrong is not incorporating side constraints into our pursuit of the good → any intuitively unjust act can be plausibly justified by utility
 - The paradigmatic example of this for Nozick is the right to life or the prohibition of killing → such constraints are imposed because the violations of such rights are wrong
- But consequentialists have 2 objections to this:
 - First, the wrongness of killing is because of the consequentialist outcome disvalue
 - Second, the badness of the consequences of multiple actions of a given sort can be aggregated so that (typically) the total badness of the results of a bunch of wrongful actions of that sort will be greater than the badness of the results of one action of that sort.
 - For example, the total badness associated with *W*'s (unprovoked) killings of *X*, *Y*, and *Z* will (typically) be greater than the badness associated with *A*'s (unprovoked) killing of *B*. It follows that it will (typically) be wrong for *A* to abstain from killing *B* if that is the only way to prevent *W*'s (unprovoked) killings of *X*, *Y*, and *Z*.
- But Nozick can deny the first consequentialist reason through deontic premises such as treating humans as ends in themselves which in turn denies the second
 - Are our intuitions right with regards to the killing 1 vs killing 5 situation (aka trolley situation)? What makes our inclinations different between the kidney and the trolley situation?
 - Our intuitions seem to be secure only when it is a matter of sacrificing the right to life for a greater quantity of lesser rights (such as property) → so a deontological approach of side constraints seem to fare better than utilitarianism at least in this respect
- Nozick's point is that human lives have distinct nonconsequentialist values in themselves that cannot be traded off with one another

The maximisation aspect of utilitarianism cannot accommodate distributive justice/fairness

- Aggregation problems within utilitarianism
 - Utilitarianism is distribution insensitive
 - Does not matter that 1 person has 10000 the utility in the world and 99 people have nothing vs 100 people have 100 utility each since both have equal overall utility

- But we want to ground distributive justice upon moral foundations; and it seems that distributive justice, in wanting to say that the distribution of utility among people matters, conflicts with utilitarianism's claim that distribution does not matter

Responses to criticisms

On Distributive Justice

- Prioritarianism: Weighted utility maximisation (betrays pure utilitarianism but better than nothing)
 - The idea is that we assign weights according to how worse off the individual is
 - This is so that in maximising weighted utility, the worst off would be prioritised even in utility space
 - Note that the weights do not affect how much utility the individual actually gets
 - But the result is that in maximising utility, it would be likely the case that the worst off are prioritised in distribution
 - Could be a democratic exercise on how to assign these weights
 - Strictly speaking, prioritarianism does not assign intrinsic value to relations between individuals, but only to the absolute level of utility experienced
 - So if one experiences an absolute level of low utility, he will be assigned a high weight → this does not change if other people are better or worse off than him
 - Though we can only determine what constitutes "low" utility by comparison?
 - Concern: on what basis do we give weights to the worse off?
 - It does not seem to arise from the utilitarian foundations because we seem to be saying that some people's value intrinsically matters more than others
 - The reasoning in assigning weights seem to be social contractarian or deontological rather than on the basis of outcome values → after all, these weights do not affect the actual amount of utility experienced
 - Does not cohere with a strong conception of impartiality: that everyone's value matters equally, *ceteris paribus*
 - Resolution on impartiality
 - While it may fail to capture the strong form of impartiality, it still captures the weak form of impartiality by maintaining agent-neutrality
 - Maintaining agent-neutrality: it doesn't matter who is the worst off
 - What matters is the rank of the individual in how much utility they have, not the individual specifically
 - We are not invoking specific characteristics of a particular individual such as causal reasons for why he is lower ranking

(such as starting out poor, or being lazy to find a job) which we may deem as irrelevant considerations

- So this weighted utility maximisation remains blind to particular characteristics of each individual
- Only concerned with abstract well-being values assigned to abstract individual lives
 - No assignment or bias of weights to specific lives → assignment depends solely on whether one has low or high utility
- Resolution on coherence with utilitarianism
 - This solution seems to be a hybrid between utilitarianism and deontology since it partially invokes deontological/contractarian reasoning in assigning weights
 - However, it still remains dominantly consequentialist since what determines rightness still depends on the maximisation of utility. Only now it is the weighted maximisation of utility
 - So rightness of an action is still dependent on outcome values, but the aggregation of these outcome values depends on contractarian reasoning
- Advantages: avoids the symmetrical pitfalls of tyranny of anti-aggregation and tyranny of aggregation
 - Tyranny of aggregation is that it does not matter that a class of people receives 0 utility, as long as utility is maximised
 - This solution makes it much more difficult for this to be the case
 - Tyranny of anti-aggregation is that we have to give up large increases in utility for marginal increases in the utility of the worst off (Rawls' leximin/maximin strategy)
 - This solution does not give rise to a lexicographical prioritisation of the worst off → a large enough increase in utility overall can in fact be better than marginal increases in utility of the worst off
- Equality as a consequentialist value
 - Expand the desirable end to include equality; basically OLT
 - Consequence
 - Then we would have to accept that a world would be better off in one respect when everyone is actually worse off
 - Consider World A where utility over 2 people is {19,1} and World B where utility is {-10,-10}

- One would be committed to saying that world B is better than world A in at least one respect (equality) despite everyone being worse off in actuality
 - Difficult to justify the intrinsic value of equality/disvalue of inequality
 - Intrinsic: implies that you are worse off just because of **the existence of someone** who is better off than you → not simply that you are concerned by this inequality
 - But utilitarianism would object to you being worse off without this affecting your subjective states
- Achieve the Nozickian/procedural notion of distributive justice
 - Empirical argument: the best world with maximal utility would be one that regulated the transfer and acquisition of goods according to Nozick's principles of justice
 - If we did not have such fair rules of transfer and acquisition, then theft and conquest would run rampant → chaos akin to the state of nature would ensue, and societal security, as Mill cites, is extremely utility conducive
 - Say that this fulfils distributive justice according to Nozick
 - Counterpoint: utilitarianism still cannot justify the deontological application of such rules
 - We can imagine a maximal utility world where Nozick's fair acquisition and fair transfer without coercion may be applied in 99 cases and not in 1 (because the utility gain from the theft is much more than the consequences of that theft → imagine it's unknown such that no one will be prompted to steal)
 - But Nozick demands that the rules be applied to 100 of the cases since they are to be applied deontologically
 - Empirical arguments are limited to wishful thinking

On Liberty Rights

- Brink's MLAU: Adoption of liberty-preserving rules
 - Follows Mill's use of rules as secondary principles
 - Mill argues that adopting liberty protecting rules would be maximally conducive to utility in the long run because it preserves societal security and societal security provides a lot of utility overall
 - So a utilitarian argument to the secondary rules but a deontic application of these rules
 - These secondary principles are only endorsed if they maximally promote the good
 - In times of conflict between rules, Mill suggests appealing to the principle of utility directly to arbitrate between the rules
 - Objection #1: Does not achieve the theoretical inviolability of rights

- The system of rules needs to meet the counterfactual condition, according to which the system of rules would be given up should it not maximise utility
- So even if we perhaps accept that the system of liberty-preserving rules are utility maximising, these liberties are still not inviolable because they are subordinate to utility
- It remains a **theoretical possibility** that these liberties are overridden should one day a set of rules that does not protect liberties are found to be more utility conducive
- So utilitarianism does not sufficiently capture the strength of the inviolability of rights
- Of course it may be begging the question against utilitarianism, for to sufficiently capture the strength of the inviolability of rights, one may have to begin from deontological assumptions → but because the inviolability of rights and the protection of our autonomy seems inextricable from our intuitions of justice, it seems that the burden is on the moral theory to accommodate for these intuitions, rather than for us to change our own intuitions
- Objection #2: MLAU makes negligible progress
 - We cannot truly be sure that, from any theoretical argument, liberty-preserving rights are in fact the utility maximising scenario
 - In fact, it is very easy to think of sacrificing one man to power the radio station (we sacrifice his autonomy and right to life/self-determination) such that the overall utility is maximised while preserving everyone else's liberties. And to meet Mill's security requirement, no one knows of this
 - But justice requires that even the violation of a single person's liberties is impermissible and unjust
- Brink: liberties as pre-eminent goods
 - Weak interpretation: liberties preserve the promotion of higher pleasures
 - Allowing for autonomy protects the promotion of higher pleasures such as reading poetry and philosophy (after all, most tyrants destroy books upon taking power → liberties could be seen to be a necessary condition for higher pleasures)
 - Strong interpretation: liberties themselves are higher pleasures
 - The experience of freedom and living could be seen to pass Mill's desire test in that any competent judge would desire these things than anything else
 - Since liberties are lexicographically preferred to everything else, they would be deemed inviolable, as $u(L1, a1) > u(L2, a2)$ if and only if $L1 > L2$ or both $L1 = L2$ and $a1 > a2$

- Given that liberties are now the primary criterion for determining utility, we can never sacrifice liberties for the sake of some gain through other pleasures
- So liberties are discontinuously better than other pleasures
- Objection: Self-defeating since lexicographical preferences cannot be represented over a real-valued utility function
 - Intuition of the proof is that by the density of the rationals in the reals, for some arbitrary value $L1$, we will find a rational number $q(L1) > q(L2)$ for some $u(L1, a1) > q(L1) > u(L1, a2)$ when $a1 > a2$ and repeat the same logic for $L2$ where $L1$ does not equal $L2$. What this function does is to map out a one to one function from the reals to the rationals which is impossible following Cantor's diagonalisation proof since any subset of rationals are countably infinite while any subset of real numbers are uncountably infinite. So there does not exist a utility representation for lexicographical preference ordering
 - So if we cannot translate lexicographical preferences into utility numbers, any attempt to incorporate rights via Mill's higher/lower pleasure distinction and lexicographical preference ordering would undermine the whole utilitarian enterprise

Misc Objections (generic anti-utilitarian objections)

Foot, Philippa. 1967/2002. "The Problem of Abortion and the Doctrine of the Double Effect." In *Virtues and Vices: and other essays in moral philosophy*. Oxford: Oxford University Press.

- Doctrine of double effect
 - Distinction between what a man foresees as a result of his voluntary action and what, in the strict sense, he intends
 - Intention: things he aims as ends and those things he aims as means to an end
 - Merely foresees: things he does not aim as ends or means as ends (oblique intention)
 - Permissible to bring out by oblique intention what one may not directly intend
 - Consequentialism is blind to this
 - Application to abortion: abortion (craniotomy) is impermissible even to save a mother's life, because to kill a child is a direct intention while the mother's death is foreseen but not intended
 - Does not mean the we do not count foreseen and not directly intended consequences (merchants who sold poisonous oil (did not directly intend deaths) and gravediggers who sold the oil to get more graves (direct intention) are both considered murderers)

- Thesis: sometimes it makes a difference to the permissibility of an action involving harm to others that this harm, although foreseen, is not part of the agent's direct intention
- Helps with some moral dilemmas
 - Case A: Judge can frame an innocent man for a murder to placate rioters (because he doesn't know who the murderer is) which would avert many deaths
 - Case B: Trolley problem
 - Foot: DDA helps us make sense of why in (B), we may kill one for five but in (A) we might not → it is one thing to steer towards someone foreseeing that you will kill him and another to aim at his death as part of your plan and in real life it would hardly ever be certain that the man on the narrow track would be killed but the judge needs the death of the innocent man for his (good) purpose
- Objections
 - Hart: why can't we take the child's death as the foreseen but not intended consequence?
 - Demonstrates the doctrine's absurdity: where is the line to be drawn between the event of the child merely "dying" and the child being "killed" or the fat man being "killed" and "die"
- Foot's better alternative: distinction is between what we do (analog: direct intention) and what we allow (analog: oblique intention)
 - Allowing has two cases: (i) forbearing to prevent (already being in the train and something that an agent could do to intervene but does not do so e.g. warn someone but allows him to walk into a trap), (ii) enabling (removal of some obstacle that holds back a train of events i.e. opening a door and allowing an animal to get out)
 - First kind of allowing (that Foot focuses on) requires an omission but there is no correlation between omission and allowing, commission and bringing about or doing
 - An actor who fails to turn up for a performance will generally spoil it rather than allow it to be spoiled
 - Sometimes allowing and causing will be seen as the same (letting relatives die of starvation) while sometimes there is a moral distinction (allowing people to starve to death in Africa vs sending them poisoned food)
 - Salmond: a positive right corresponds to a positive duty, and is a right that he on whom the duty lies shall do some positive act on behalf of the person entitled. A negative right corresponds to a negative duty, and is a right that the person bound shall refrain from some act which would operate to the prejudice of the person entitled. The former is a right to be positively benefited; the latter is merely a right not to be harmed

- Defective as general account of rights but we can think of negative duties as obligation to refrain from killing/robbery etc and positive duties as to look after children or aged parents (and extended to charity)
- Application
 - Trolley problem: weighing one man's negative right against five men's negative rights (conflict of negative duties) and thus should just do the least harm he can
 - Judge problem: has a positive duty of bringing aid to the innocent people threatened by rioters but a negative duty of not harming an innocent man
 - Negative duty can be said to be lexicographically prior to positive duty
 - Abortion I (life of mother can be saved by killing the child and vice versa): conflict of negative rights and possibly equal so no strict conclusion
 - Abortion II (life of mother must be saved by killing the child, but if nothing is done, the mother will perish and the child can be safely delivered after the death): DDA says that we may not intervene and the positive/negative rights distinction says that it is a positive right against a negative right (especially in the case where the child is born already) but Foot highlights that some people think if enough people's death were involved, it would be permissible to sacrifice the one, and people view the mother's life vs the unborn child as akin to a many-to-one comparison
 - Is this more like the judge's case or the trolley case (but 1-1 so indeterminate)
- Positive and negative duties are the main generator of our intuitions: direct and oblique intention only play a subsidiary role
 - But also important within positive duties whether it is a strict duty or an act of charity (i.e. helping my children vs those far away)

Part II: Forms of Consequentialism

Act vs Rule Consequentialism

Conformity Rule Utilitarianism

- Conformity RU imagines a world where everyone follows the established system of rules
 - Not only is it costless is helping people understand the rules, but it also assumes restlessness in people carrying out the rules
 - Assumes that perhaps people are machines dutifully carrying out the rules
- Problems
 - Collapses into Act utilitarianism
 - We cannot internalise the cost of the decision-making procedure and therefore the optimal decision procedure would be to add as many exception clauses as possible (do not lie unless it prevents a disaster) until it reaches “do not lie unless lying produces greater utility” which essentially means, do the action which produces the greatest utility
 - This is because conformity RU ignores the costs of people having to learn, understand and implement the set of rules in their daily lives
 - There are some rules (such as retaliation rules) where the world would be better off if not complied with
 - The rule that prescribes that you retaliate against attackers: it is optimal for people not to ever have to comply with the rule by not attacking one another in the first place → so there are empty rules that are best not to be complied with
 - With 100% compliance, there would be no place for rules of punishment
 - Rule consequentialism cannot generate or justify rules about how to deter murder, rape... by unmitigated amoralists, unless rule-consequentialism picks its rules with reference to an imagine world where there is internalisation of the envisaged rules by <100% of the population
 - Morality cannot allow for me to bear the cost of the rule by following it and let others benefit from it by not complying with it yet reaping the benefits of general compliance
 - <100% compliance (acceptance) allows for contingencies in rules: Imagine 2 children starving and the rule is you and your friend are obliged to save one each, but your friend does not save him
 - Imagining an imperfect world gives us the space to add in the exception “if others don’t help, do more than what is required of you to prevent the disaster”
 - We could reach an optimific rule: People should help others in great need when they could do so at a modest cost to themselves, cost being assessed aggregately, not iteratively (Cullity 1995)

- Respects the cumulative nature of the cost of helping others as potentially too costly if iteratively assessed

Acceptance Rule Utilitarianism

- Acceptance RU imagines a world where everyone accepts these rules but does not necessarily comply.
 - It does not immediately assume society's seamless ability to comply, and instead begins from the point of society learning these rules.
 - Therefore, the acceptance form of RU makes explicit the rules' costs of internalisation
- It recognises that rules themselves have costs, and not just benefits from the actions they sanction.
 - As rules become more complex, they become more difficult to learn and act upon, incurring large costs onto society
 - So the optimal situation would be for a set of rules that is dissimilar to the act utilitarian decision procedure (though this is an empirical question)

Arguments for Rule Utilitarianism

- Provides a superior decision making procedure
 - Act utilitarianism does not (though this remains to be seen what is the utility maximising world once costs of decision making are internalised – an open empirical question)
 - Act Utilitarians make an explanatory object to Rule Utilitarianism by arguing that it misidentifies what makes actions right or wrong (it is not the violation of rules but the non-maximising of utility)
- Contractualist premise rather than a consequentialist one
 - Contractualist premise is that rules are morally justified just if they are ones to which everyone has good nonmoral reasons to agree
 - The argument then proposes that an act is morally permissible if allowed by morally justified rules
 - Then the argument contents that the only rules that everyone has good nonmoral reason to agree to are the rules whose widespread acceptance would maximise expected value impartially
 - This may not be the case (Rawls argues that people would adopt the maximin principle in such a contractual scenario instead)
 - 'Good nonmoral reason' is used in the sense of maximising welfare
- Taps on intuitively plausible ideas
 - What if everyone felt free to do what you are doing → contractually agreed to
 - Should we not try to follow a code that could plausibly maximise welfare?
- Avoids the breaking of rules for marginal gains

- Act utilitarianism would allow for a breaking of promises even if the utility gain is marginal → counterintuitive
 - But we would also believe that breaking promises for a large gain in utility would be justified
- Act utilitarianism and absolutism (the view that some acts are absolutely right or wrong with no exception) are on opposite ends of the spectrums since AU proposes that actions are only conditionally wrong
- Rule utilitarianism advantage is that it coheres with when we intuitively feel we can and cannot break rules for the sake of the overall good
 - This is due to the fact that we can only have disaster prevention exception clauses and not marginal gain clauses → general internalisation would be optimal
- Perhaps rule utilitarianism is self-standing insofar as it only refers to the utility principles as its justification and meta-selection principle and not in its decision procedure and its criterion of rightness
- Avoids the demandingness objection
 - Act Utilitarianism would require us to sacrifice much of ourselves for others → our singular enjoyment <<<< than others suffering
 - But morally unreasonable to demand such level of self-sacrifice → perpetual self-improvement > than what morality demands → perhaps imply supererogatory
 - Imagine allotting offices → you are a senior and you get a corner office but someone benefits more than you and you are indifferent. Even if you do not give up your office for him, no one will say you are morally condemnable → because he has no right to your office in the first place
 - Morality does not require us to sacrifice ourselves normally
 - No space for personal prerogatives
 - But note that RU may still end up being extremely demanding – we just do not know
- Avoids too permissive objection (continuous with the integrity objection)
 - Rule-utilitarian would argue that the best system of rules will incorporate important constraints
 - Would not permit doctors to cut up healthy patients and distribute their organs to save others → because in the aggregate, terrible consequences
 - But this remains an open question
- Motive Utilitarianism: Act Utilitarianism's decision procedure is not simply suboptimal in committing us to overly time consuming welfare calculations, but also **provides us with suboptimal motivation structures which subtract from our overall welfare** (Robert M. Adams)

- Adams proposes that the motive of the the action itself contributes to the outcome of the action in affecting our psychological states
 - Note that this can extend to any form of utilitarianism: pleasure, desire, OLT
 - Appreciating beauty may be a value in the OLT, we may gain more pleasure from intentionally understanding the paintings, and it may be our desire to revel in art
- I think Adams' point is this: that **adopting the motive of maximising utility** as per AU **does not provide us with the utility maximising motive structure**, because it takes us away from enjoying the activities we care about or could care about
 - Example: Jack is studying sculpture on the stone choir screen → enjoying this less than other parts of the cathedral and will not remember it well → more happiness on balance if he left earlier
 - This action would not be advised by the AU decision procedure and may be considered wrong if we do not internalise the utility effects of the motive
 - But it may be the case that his caring more about seeing the cathedral than about maximising utility has augmented utility, through enhancing his enjoyment, by more than it has diminished utility through leading him to spend too much time at Chartres → thus he did right by motive utilitarian standards → the motive enhanced his enjoyment because he was not looking at the cathedral for some instrumental end but rather for its intrinsic utility and thus could appreciate it better
 - He may desire to enjoy the paintings; but if he is only focused on maximally fulfilling his desires, would **this not conflict with his ability to actually enjoy the art around him?** And for pleasure, would the **perpetual concern** for utility maximisation **not subtract away from our overall pleasure?**
- Utilitarianism then cannot simply evaluate utility on the basis of **what action** is committed, but must internalise the utility effects of **how** the action is committed
- So Adams extends the point that AU's decision making is not utility maximising: it is not utility maximising in also how we go about conducting our lives → being rational calculators of welfare itself undermines our ability to experience pleasure (in the things around us) or truly satisfy our desires
 - This is beyond the objection that it consumes too much time and thus cost us much
- So in internalising the utility effects of the motivation structure, Adams concludes that the AU motivation structure is not utility maximising

- We can see how RU can plausibly **incorporate having the appropriate motivation structure as a rule**
- This shows that RU provides greater flexibility in our decision procedures than AU
 - Disagree with Adams on the incompatibility between motive utilitarianism and AU: motive utilitarianism is not incompatible with AU's explanatory component/criterion of rightness of what actions are right – we just need to extend the analysis of actions to the dimension of how they are committed/their motive since it has utility effects
- The second key insight from Adam's is that RU's decision procedure is not necessarily utility maximising because it retains an inappropriate motivation structure; so **MLAU does not necessitate rules as the decision procedure**
 - RU prompts people to be **motivated to act in a certain way because of the established rule** and not for the sake of the action itself
 - Motive utilitarianism requires Jack to adopt a motivation structure that inherently appreciates the value of revelling in the art around him; it is a motivational structure that compels an agent to act for action X because **he appreciates the value of action X**
 - The extreme reading is that any motivational structure other than acting for some action X itself incurs some costs in the decision procedure (of course the intensity varies); so we need to internalise the cost of the suboptimal motivation structure: makes the question of whether MLAU incorporating RU as decision procedure is max(u) even more of an an open question
- Congruent with the universalisation test
 - If your friend tells you that it is fine for him to evade taxes or cheat in exams because the deviation from generally upheld norms will not cause any bad consequences, your reaction will likely be 'What if everyone thought it would be OK to do that?'
 - The moral rules are universally applied and upheld → static in nature
 - The moral imperative does not change based on what other people do → when they don't help someone, it does not force you to help more so it avoids the too demanding charge
 - But we also find it counterintuitive if our moral rules are constantly in flux

Objections to Rule Utilitarianism

- Criterion of rightness incoherent with the principle of utility

- Principle of utility: about the individually utility maximising action → so RU deviates from this completely by pegging rightness to conforming to the set of rules
- Solution: suggest that the rule-based criterion of rightness still explains rightness and wrongness on the basis of maximising utility, and that it is actually more coherent with achieving a history/world with the greatest utility since it can account for the **interaction of actions across time and between people**
 - Crisp's edit of the principle of utility: The best actions are those which feature in the best possible history
 - Decompose the notion of history: not just about actions across time, but between people as well
 - So, the principle of utility sanctions a set of actions that collectively maximises overall utility, rather than individually utility-maximising actions → AU's decision criteria commits the fallacy of composition (see W3 essay and extend the argument to across time)
- Revised Objection: even if we accept Crisp's suggestion, RU assumes that the in set of actions that collectively maximises utility would include rules of actions that apply generally
 - RU assumes that, usually, the set of actions that collectively $\max(u)$ implies that we must always not lie since we create a general rule dictating whether or not it is acceptable to do an action
 - But intuitively, the maximising set of actions would have deviations from the rule across time and between people (in some cases, it would be best not to lie and in others it would be best to lie)
 - Even incorporating contingencies may be insufficient; in speaking about the criterion of rightness, we can ignore the costs of decision making
 - So it is hard to accept that the RU criterion of rightness adequately captures the principle of utility
- No unique best system of rules
 - Implicitly assumes that there would be one unique best system of rules
 - And the Utilitarian principle supposes indifferences between two actions (and by extension, sets of rules) that are maximally and thus equally conducive of welfare
 - But this indifference cannot hold since everyone needs to adopt a single system of rules
 - Consider 2 systems that one differs only in asking you to drive left or right → not implausible to fill in the details so that these systems are equally good

- Utilitarianism here is indifferent between you following a set of rules that differ in asking you to drive left or right
 - So people drive on both sides simultaneously and crash and de
- But this objection is mistaken → rule utilitarianism would **only be indifferent between the set of rules at the point of choosing the set of rules**
 - It would not allow for one to implement both set of rules simultaneously because it is obvious that such simultaneous implementation would be mistaken
- Assumed existence of a best set of rules distinct from AU
 - Remains an open empirical question → not a denial of RU, but cautioning how much progress RU actually makes over AU
 - It may be the case that even after internalising the cost of decision procedures, AU's decision procedure is still the best
- Problem with partial compliance/acceptance
 - Require rules that handle people who transgress rules even though majority accept the rules (and some do not comply)
 - Idealistic: If everyone accepted my pacifist principles, there would be no wars and no need to take up arms. Therefore, a rule about not taking up arms should be established
 - Non-problem: obviously just incorporate rules on punishment conducive to the overall welfare → compel people to comply too
- Disaster prevention
- Incoherence with common sense in why actions are wrong
 - Fundamental disagreement about what makes actions right or wrong (justification)
 - Rule Utilitarianism: what makes actions wrong is that fact that this action is forbidden by a system of rules whose general acceptance would bring about the greatest total of well-being
 - Common sense: what makes it wrong is something about you and the person you made a promise to is simply the fact that you made a promise to him
 - Awkward: people do not follow rules for the sake of following rules, but for some underlying reason → so it attacks RU's criterion of rightness for assessing rightness from the standpoint of following rules itself
 - Solution: suggest that the rule-based criterion of rightness still explains rightness and wrongness on the basis of maximising utility, and that it is actually more coherent with achieving a history/world with the greatest utility since it can account for the **interaction of actions across time and between people**

The Collapse Argument (2 strands)

- First strand: Rule Utilitarianism collapses into Act Utilitarianism
 - If what is morally right is the action that produces the maximally good outcome, then the set of rules would eventually have to specify infinite contingencies until the decision procedure collapses into “Do X unless it does not maximise the good”
 - But obviously we can tackle this using acceptance Rule Utilitarianism
- Second strand: Rule Utilitarianism collapses into customary morality
 - If you do not include exceptions, then you are not a consequentialist because you do not truly use utility/outcome values in your analysis
 - Following rules are not necessarily consequentialist; the justification of these rules must be consequentialist (that they produce the greatest utility) → so RU threatens to depart from utilitarianism itself in its explanatory aspect
 - Solution: we will probably have to include some exceptions
 - Our utility maximisation problem would be to maximise utility with respect to the complexity of the rules such that the net utility from “utility from a world following the rules - the cost of internalisation” is maximised
 - This implies that it is not necessarily the case that we have simple rules of thumb as in customary morality
 - RU is still grounded on maximising utility – it just needs to take into account the costs of having complex rules
 - Moreover, by needing to consult the Act Utilitarian principle (which is itself a rule) during times of conflicting rules, it shows that the fundamental principle is grounded on utility and maximising outcome values
- Brad Hooker’s conjecture: we can cleanly avoid the collapse argument by not beginning from an overriding commitment to maximise utility
- Bykvist Objection of the Collapse argument

- Bykvist asks us to imagine a simultaneous single-shot game

A/B	B go to Oz	B go to Kings
A go to Oz	6	0
A go to Kings	0	10

- If we both go to Oz, that would be right according to act utilitarianism
 - This is because act utilitarianism takes the other’s action (going to Oz) as a given – exogenously determined
 - So AU evaluates an agent’s action on whether it maximises utility in a given scenario, even when the factors are beyond the agent’s epistemic standpoint

- And we want to say they are both wrong! (but Bacceli contests this)
- The point is that acting according to act utilitarianism will not produce the best circumstances in itself → but this is not acting according to act utilitarianism → the evaluation is ex post
 - RU on the other hand would imagine the set of actions that **collectively determine** the optimal outcome (does not take any action as exogenously determined since RU follows a universalisation test)
- Perhaps we need to redefine the best history as one that is in relation only to the current agent's actions (so subject to exogenously determined actions)
 - After all, the right action intuitively is influenced by other player's actions → like donating, or saving the child

(Further) Sophisticated Consequentialism

Fully informed desire (ideal observer) accounts

- The life that would be best for a person is one that they would desire if fully informed of the facts about various options available to them, and their greatest welfare would consist in the fulfilment of that preference or desire
 - So someone who has been brought up to despise happiness (Miserable Maud) would realise that she had not been presented with accurate facts and having been presented with all relevant facts, would desire a life of happiness
 - How can we tell if someone has been fully informed, if not by acting in an expected way, which in turn means that we presuppose people to desire what is independently/intuitively considered as good

Objections against the Fully Informed Desire account

- Narrowness objection
 - There seem to be things that are good for us independent of our desire for it
 - Curing cancer for someone who does not desire the cure seems to be good for the person regardless of his desire
 - But we can bite the bullet: perhaps when we say there are things that seem to be good for us independent of our desire for it, we draw on evidence such as curing cancer for someone who hates it → but what we are drawing on is our own subjective conception of the good; we assume that just because curing cancer is desired (and thus good for) many people, it follows that it is good for that person too
 - But this doesn't follow → if we begin from a premise of subjectivity, the fact that we all think that curing cancer is good is a coincidence or is a result of some causal process irrelevant to the good
- Even fully informed desires can still lead to rather **trivial and meaningless desires**
 - The grass counter desiring to count the number of blades of grass in the garden → could we say that he is truly better off?

- The fully informed account wants us to avoid having bad desires; but if it dictates what desires are good, this surely implies that there is some pre-existing criterion of goodness that underlies desire
- Dilemma: desire accounts must either have some **independent criteria** of what constitutes a 'good' desire or must be simply **held hostage to contingency**
 - Horn #1: Desire accounts argue that desires are not open to rational criticism in light of what is good because **there are no independent facts** about what is good
 - What someone desires cannot be criticised for being less good for them than some alternative other than in light of the desires that person actually has, or would have in certain conditions.
 - Desire accounts therefore leave its theory of value hostage to contingency
 - One can bite the bullet that what is good is necessarily relative to the person → it is necessarily contingent and that does not matter because fulfilling one's desire would be seen as valuable to that person
 - How can we criticise another's desires? When we do, what we are actually saying is that given my motivational structures and upbringing/background, I would not desire what he desires. Given that one's desires are likely causally related to psychological and upbringing factors, one's could be said to be defensible according to his traits
 - Horn #2: To avoid contingency, desire accounts seem committed to some **independent criterion of the good**
 - Again, how can we tell if someone has been fully informed of the facts since it requires us to demonstrate that we understand the importance and relevance of these facts → such demonstration requires us to act in a particular way
 - This implies that we import pre-theoretic substantive judgements on what is considered good → so we desire something because it is good rather than something is good because we desire it (fails the explanatory account of the desire theory)
- Disinterested desires
 - Imagine Sagan who desires that we make contact with aliens and we manage to do so in 200 years. Unfortunately, Sagan does not live that long and thus is unaware of this → can we say he is better off? No!
 - I think this thought experiment is unfair. The reason why we think that Sagan is not better off is not simply because he is unaware of it but because he is dead. A mother who wishes for her lost child to be loved would still be better off if the child receives love without her knowing it

- Rosati's objection against fully informed desire accounts
 - Essentially argues that any account of what it is like to have "full information" of desires is impossible for two reasons
 - **Irrelevance Objection:** Any judgement arising from such a standpoint would be normatively vacuous → the fully informed person, though purportedly you, may not be someone whose judgements you would recognise as authoritative, since the ideal advisor is too far removed from who you are as a person
 - **Impossibility Objection:** Because of what it is like to be a person and to have a perspective, it appears that no one can be truly fully informed
 - Rosati's requirements of a theory of ideal well-being: Railton argues that we must be sufficiently engaged with what is good to us; an ideal needs to move us to aspire it rather than simply having a subjective account → one must hold not only that an adequate account of a person's good must affect a motivational connection between a person and her good, but it must also satisfy a double motivational link
 - First link: effects a motivational connection between the individual and her good under "ideal conditions": X can be good for a person only if she would desire it for her actual self under ideal conditions
 - Ideal condition requirements seek to avoid the wholly contingent nature of desire preference → something is not good for someone simply because she desires it
 - Second link: effects a motivational connection between an individual and information about her counterfactual desires or concerns
 - Conditions can only be relevant if the actual person would care about the fact that something would prompt her concern *under those conditions*
 - What this means is that the conditions in which the "idealised" me desiring X motivates me to actually desire X cannot be too far removed from reality → otherwise, why would I care if that "me" who desires X (though entirely objective) is not realistic at all?
 - The first condition is about fully informed desires as being relevant to the agent (she would actually desire it under ideal conditions) and the second conditions anchors these ideal conditions to be adjacent to reality
 - 2 notions of alienation:
 - First sense: one is incapable of caring in any conditions or only cares under certain conditions which are irrelevant to the concerns people

typically have when they wonder whether something is good for them (i.e. X is good for me if I am in Mars)

- An internalist conception of the good (that what is valuable to someone emerges from within that person and not from external forces) → so this account of the good consists in what she would want under conditions of sensory deprivation; it must incorporate only counterfactual conditions that bear reasonable relationship to the concerns that animate inquiry into our good — concerns we have under ordinary optimal conditions
- Second sense (though quite similar): judgements and their recommending force depend on whether this is a conceptual or merely contingent feature of such judgements
 - Unless the person cares under ordinary optimal conditions, an account of a person's good would be hard pressed to explain how it is that judgements about her good normally have recommending force
- So desires can be more defensible to the extent that they are better informed and rationally based → but it cannot be under extreme conditions such that they arise more from the conditions the agent is situated within than the agent's rational desires themselves
- Understanding the Ideal Advisor/Fully informed rational desire account
 - **Objective 1:** Seeks to avoid importing any substantive evaluative judgements into their accounts
 - Otherwise, this compromises the explanatory component of desire theory → that something is good because someone desires it (creates an independent criterion)
 - **Objective 2:** Isolate with respect to each person what someone like her would come to desire with full information → allows for each individual to serve as her own ideal advisor
 - So this is contingent on her innate preferences (internalism)
 - But if desires are then a function of information, how can we know that someone is fully informed other than seeing that changes occur in her desires/preferences (which defeats Objective 1)
 - **Objective 3:** X should move you under specified ideal conditions and will have recommending force even if X does not actually move you at current
- Notion of being fully informed (vague)

- Brandt: “ideally vivid way” → “gets the information at the focus on attention, with maximal vividness and detail, and with no hesitation or doubt about its truth”
- Railton: vividness depends upon a person’s receptivity which in turn depends on education or experiences that can animate this vividness → the person must be able to appreciate the *importance* and *relevance* of the facts provided
 - Again, our definition or test of when a person “appreciates” information likely presupposes pre-theoretic judgements because we want to see one’s desires change in a certain way
 - Imagine a sadist who desires people to suffer → we provide him with facts about how this actually harms him and harms other people (note that this already presupposes his desires are bad – can facts even be neutral?). If the sadist keeps his sadistic desires, we would think that he has not appreciated the facts! Our measure of appreciating the facts is calibrated to our pre-theoretic substantive judgements – this compromises the explanatory component of desire theory
- Rosati tries to salvage the interpretation: she interprets vivid in the sense that information has been rendered maximally vivid when its motivational impact (indicated by her reactions) as a function of types of representations has been exhausted (this is still vague)
 - Avoids the hopeless task of trying to explain which representations of information are most vivid
 - Avoids treating information as maximally vivid only when it motivates the recipient in a specific way
- **Impossibility Objection:** no one can be fully informed
 - How we assess any experience is contingent upon our prior and subsequent experiences as well as our traits and motivational systems
 - This leads to the problem of **experiential ordering**: Poverty after wealth is experienced differently from poverty after poverty or wealth after poverty
 - We often re-estimate the value of X after experiencing something new → so this implies that the process of fully informing a person will need in some way to offset the effects of experiential ordering, by requiring that a person experience all lives in all possible orders
 - **Impossibility #1 (unrealistic):** Requires immense cognitive ability that is clearly impossible
 - **Impossibility #2 (self-defeating):** One cannot compare the different infinite lives to assess which is rationally desirable (because one’s

assessment of the life is based on fundamentally different traits and motivational structures shaped by the different experiences themselves) → incommensurable

- One's first experience as a possible life as a Marxist atheist is grounded upon atheistic traits that are developed throughout that life vs one's experience as a nun would be grounded upon traits like comfort in the enigmatic
- In comparing these lives, she would only find the life of an atheist desirable if she had the relevant traits and dismiss the life of the nun as undesirable and vice versa
- **Our features vary as we progress through life, so our features are not the same across lives and across time within those lives**
 - Significantly, we are not the same persons through the one life we actually lead AND across the lives we imagine we lead → if we are made who we are by our features → our features are a by-product of the processes she must undergo in order to be fully informed
 - So there is no such thing for "what the experience is like for me" → "what is the experience like for me, given what I am like at time T, and in life A" → if we can't have interpersonal comparisons of preferences because of how subjective this is, this undermines the intent of the full information account
- It becomes difficult for the agent to compare between these different lives to figure out which life of desires is optimal because he is a fundamentally different person across different lives. He cannot consider "what the experience is like for me", but "what the experience is like for someone with X traits".
- Lives are evaluated along infinite dimensions {contentment, challenging, peaceful, meaningful..} → weighing these dimensions against each other depends on one's traits which are invariably affected by the experiences themselves → we reach a dilemma: either (1) we drop the requirement of needing to be fully informed by changing our psychological features so that we can evaluate experiences from a consistent standpoint (but this will likely not move one's desires) or (2) we lose our ability to rationally compare between these lives (since my experience of one life is grounded on fundamentally different characteristics in order to appreciate them fully)

- **Irrelevance Objection:** being fully informed renders the ideal advisor to lose her normative force in motivating you to desire X
 - (i) A person would have to change markedly to become fully informed;
 - (ii) Ideal Advisor views lack the resources to guarantee that the fully informed person, though purportedly oneself, is someone whose reactions an individual either will or should regard as authoritative
 - **Since one's (psychological) features affect what can be received and how information will be assimilated, full information accounts require us to go through extensive changes in order to even receive information**
 - Consider a person who lacks empathy but holds an ideal of complete honesty → she kindly declares to you her real feelings about you or assists you in realising your true faults → information about how brash she is or how her controlling activities affect those around her cannot be informing for such a person
 - What can be informing for her (and for us in general) is a function of her (our) personality structure or who we are at a particular time
 - So we likely need to change our personality structures (which make us who we are) to truly assimilate information → but yet, we need some sort of a personality structure to make any judgement about information received → there does not seem to be any "objective" or correct personality structure → there is no reason why rationality (which may be too weak to itself be a structure) should be deemed correct
 - Note that even what kinds of personality structures best appreciate facts are indeterminate: **does appreciating facts require us to have personality structures that engender a positive OR negative response?** As Rosati points out to avoid pre-theoretic judgements, it could simply be any reaction → but then we are **indifferent between 2 personality structures** and we **HAVE to import pre-theoretic substantive judgements to choose which personality structure should be adopted**
 - Moreover, given the problem of experiential ordering, the person's new desire arises from an agent radically different from the person who underwent idealisation
 - Our experiences in those lives shape our traits and motivational systems → they are not shielded from one another nor fixed in place but have a mutually affecting relationship

- So the ideal advisor becomes nothing like me and thus loses normative force
- Some responses to Rosati
 - Our traits only affect our appreciation of specific instantiations of welfare values, but not classes of welfare values
 - Distinct traits would not render us unappreciative of classes of welfare values → like pleasure
 - It would only increase or decrease our appreciation of specific instances → like sleeping peacefully at night after experiencing war
 - But it is obvious that there exists counterexamples where traits can affect our appreciation of the entire welfare value class (Miserable Maud)
 - Instead of adding features, be like Rawls and strip back features
 - Rosati tentatively considers this but it is easy to see why one rejects it → Rawls incorporates his own thin theory of the good to motivate people to react a certain way → without contingent features it is difficult to assess how one might react or if one can even sufficiently appreciate the facts
- Conclusively, Ideal Advisor theories cannot figure in a plausible analysis of our notion of a person's good
 - Problem of how people might occupy a perspective that gives her equal access to lives in which she would have conflicting traits: these perspectives given by a person's possible lives involve not only differing traits (over time) but conflicting commitments and belief systems
 - The fully informed person will have vastly different belief system from that of the persons she would be in her various possible lives → affect her efforts to call up and compare what these lives are like
 - The problem is also that some lives are conceptually so that a person cannot be simultaneously be fully informed about both, so a person cannot be fully informed
 - Some lives are so deeply in conflict even the notion of lives being possible for someone is vague
- Conclusion
 - Can neither capture the normative force of our notion of a person's good (doesn't engage the agent sufficiently) nor does it explain how a person can be fully informed (and appreciate the information)
 - Need to fix non-cognitive features but it is not obvious how to without importing substantive evaluative judgements or capture the

normative force/provide a coherent picture of the fully informed person

- Adaptive desires
 - We can simply change our desires to conform with the world to make ourselves better off → Stoic philosophy
 - Though there definitely exists deep-seated desires
 - Adaptive desires are problematic because we can change our desires to things that seem unworthy of concern
 - What makes a person better off is not simply that he gets what he would favour but also that these things are worthy of concern
 - Perhaps it is still true that nothing that is not desired by him can be good for a person → but this does not imply that everything desired by him is good for him: desire is therefore at best a **necessary condition** in the theory of value
 - The poor can get used to their state of affairs such that they become comfortable with whatever they have and desire nothing more
- Makes self-sacrifice impossible
 - If one sacrifices themselves by running into a burning building to save a child, because he desires to do so, he cannot be said to be worse off (since what is good for him is what he desires)
 - Clearly an impotent objection because he has other simultaneous desires (if not at the instantaneous point in time then intertemporally) which he gives up (such as continuing his life and pleasures)

Can everything be consequentialised?

Brown, Campbell. 2011. "Consequentialize this." *Ethics* 121 (4):749-771.

Central Claim: rejects the claim that consequentialism is empty because one can consequentialise any theory – Campbell argues that consequentialism has substantive, unique content and cannot consequentialise any theory, because consequentialism is characterised by three conditions: (i) agent neutrality, (ii) no moral dilemmas, and (iii) dominance. Campbell claims that there are some moral theories which violate each of these

- Consequentialists vary their “theory of the good” because they are not committed to any theory of the good
 - To reconcile consequentialism with the view that this action you’ve described is wrong, we need only to find an appropriate theory of the good, one according to which the consequences of this action would not be best
 - Take the features of an action that the theory considers to be relevant, and build them into the consequences; for example, if a theory says that promises are not to be broken, then we restate this requirement: that a promise has been broken is a bad consequence. We can even further include absolute side

constraints by giving the forbidden action a lexically prior negative weight to promise breaking

- Let M be any moral theory which we aim to consequentialise. Let C be consequentialism. So a particular consequentialist theory is equivalent to some conjunction C&G, where G is a theory of the good. Say that the deontic output of a theory P is the set of all propositions about rightness of wrongness of actions which are entailed by P; and say that P and Q are deontically equivalent iff they have the same deontic output. Then, to consequentialise M, we simply find some theory of good G such that M is deontically equivalent to the particular consequentialist theory C&G. We might then call C&G a consequentialist counterpart of M
- This strategy is self-undermining for the consequentialist: consequentialism becomes a triviality
 - Consequentialist–Alternative duality: There exists a consequentialist theory C is the dual of some alternative theory T, for all Ts
 - Campbell thinks that this speculation is misplaced → even if we make consequentialisation as easy as possible, he shows that the project of consequentialising all theories cannot succeed
- Maximising the good: to makes things as good as possible, the best they can be → non-negotiable for consequentialism
 - Potential procedure: (a) partition the world into a set of mutually exclusive and jointly exhaustive things (to avoid double counting), (b) then assign to each thing in the partition a number representing the degree to which it is good, (c) add up all the numbers and this sum will represent the extent to which things are good overall
 - Controversial: the goodness of a whole is uniquely determined by the goodness of its parts → a state s_1 may be x good and a state s_2 may be y good, but state s_1, s_2 together is $z > x+y$
 - General picture: when a person acts, she faces a range of alternatives, different actions she might perform. For each action, there's a unique possible world that would obtain if the action were performed → this is an outcome of the action → faces a choice between possible worlds
 - Action is right iff its outcome is best
- Consequentialism is understood as a criterion of rightness rather than a decision procedure
 - Can have different decision procedures but a consequentialism theory of rightness
- Consequentialism is understood as a purely extensional thesis
 - Being morally right and maximising the good are necessarily coinstantiated → makes no claim of causation, determination or explanation → need not say that the acts are right because their outcomes are best

- Consequentialism is an objective theory
 - Since the outcome of an action is some possible world, a person cannot in general know what the outcomes of her action will be → what is considered are not someone's beliefs or subjective probabilities, but objective rightness
 - Depends on what in fact the outcomes of actions are
- Theories of the good – consequentialism is not a moral theory until combined with a specific theory of the good
 - Utilitarianism: hedonistic theory of the good
 - Railton's Objective List Consequentialism: pluralistic theory of the good – includes several goods which are viewed as intrinsically nonmorally valuable such as happiness, knowledge, purposeful activity, autonomy, solidarity, respect and beauty
- Prima facie Objection: does constructing a consequentialist counterpart for M amount to the same thing as consequentialising it?
 - Suppose I employ you to paint my red house blue and you proceed to build another house which is a perfect replica of mine except it is painted blue rather than red – then you haven't done your job. Likewise, C&G may be consequentialist, but that is a different theory since C&G is not M, even if they agree on all deontic questions, about which actions are right or wrong but they need not agree on questions of the good, because M need not imply G, the theory of the good implied by C&G
 - Reply: consequentialiser needs to adopt a semantic thesis about the meaning of the term "the good" as a theoretical term → a term which is introduced by a theory as a name for something whose existence the theory postulates (like theory of phlogiston postulating the existence of phlogiston). The consequentialiser treats 'the good' as a theoretical term introduced by consequentialism, in the same way that 'phlogiston' was introduced by phlogiston theory – that is 'the good' names the thing, if there is one, which has the properties attributed to the good by consequentialism
 - The good is thus, by definition, that which ought to be maximised or more precisely, that thing the maximisation of which is both necessary and sufficient for acting rightly
 - Imperfect analogy: phlogiston is not used outside the theory, while 'the good' is used outside consequentialism → but we can think of other cases where an ordinary word is co-opted, and the case here is analogous
 - Hence, let 'the good G' denote whatever is good according to G. As we know, an action is right according to M iff it is right according to C&G. And an action is right according to C&G iff it maximises the good G. Therefore, an action is right according to M iff it maximises the good G. That is, M implies that the

good G (and nothing else) ought to be maximised. Given our semantic thesis, M thus implies that the good G is the good

- Motivating a counter-response
 - The article “The” is crucial: “the good is that which ought to be maximised” is not equivalent to “Anything that ought to be maximised ought to be maximised” which is a tautology
 - The sentence of the form “The F is G” is true iff the definite description “The F” denotes something which is G while “Anything that ought to be maximised...” can be vacuously true even if nothing ought to be maximised
 - Furthermore, a thing is denoted by “The F” iff it, and only it, is F $\rightarrow$ “The F is F” is true iff exactly one thing is F $\rightarrow$ so “The thing which ought to be maximised ought to be maximised” is true iff exactly one thing ought to be maximised
 - Thus, there remains a way for a moral theory to resist consequentialisation: by implying that nothing ought to be maximised (and more than one thing ought to be maximised)
- Limits of consequentialism: Formal Model
 - Primitives: possible worlds (complete ways the world might be) and agents (creatures capable of actions apt for moral evaluation)
 - A choice situation is an ordered pair $\langle i, A \rangle$ where i is an agent and A is a finite, non-empty set of worlds
 - These worlds are the alternatives between which the agent i must choose, the outcomes of the various actions which she might perform in this situation
 - Assumption: every such ordered pair constitutes a choice situation
 - Deontic output of a moral theory will be represented by a rightness function R which assigns to each choice situation $\langle i, A \rangle$ a subset of A , denoted $R_i(A)$
 - Worlds in $R_i(A)$ are the alternatives which are right (permitted) in the situation $\langle i, A \rangle$, according to the moral theory represented; the worlds $A \setminus R_i(A)$, that is, those in A but not in $R_i(A)$, are the alternatives which are wrong
 - A theory of the good will be represented by a complete order $\geq$ on the set of all worlds
 - Reflexive, Transitive, Complete/strongly connected (linear pre-order)
 - If $\geq$ is an order on S , and X is a subset of S , then we may define the set of ‘greatest elements’ in X , relative to $\geq$, as the set:
 - $\text{grt}_{\geq}(X) = \{x: x \in X \text{ and for all } y, \text{ if } y \in X, \text{ then } y \geq x\}$
 - That is, the greatest elements are those which are at least as great as every element

- We may then say a rightness function R 'maximises' a world order $\geq$ iff in every choice situation, R selects all and only the greatest alternatives relative to $\geq$; that is, for all i and A , $R_i(A) = \text{grt}_{\geq}(A)$
- Consequentialism: R maximises exactly one complete world order $\geq$
- A moral theory is consequentialistic iff it is represented by a rightness function R which satisfies C
 - Proposition: R maximises no more than one world order
 - Proof: suppose $x \geq_1 y$ and $\neg x \geq_2 y$. Then R maximises $\geq_1$ only if $y \in R_i(\{x, y\})$ and R maximises $\geq_2$ only if $y \notin R_i(\{x, y\})$. So, R cannot maximise both $\geq_1$ and $\geq_2$
 - Corollary: by C , R maximises at least one world order
- Thesis: there are some rightness functions which do not satisfy this condition of representation i.e. correspond to moral theories which cannot be consequentialised
 - Want to show that violates three conditions
 - Agent Neutrality (AN): $R_i(A) = R_j(A)$
 - No Moral Dilemmas (NMD). $R_i(A) \neq \emptyset$
 - Dominance (D). Suppose $\{x, y\} \subset A \cap B$, $x \in R_i(A)$, and $y \notin R_i(A)$. Then $y \notin R_i(B)$
 - Theorem: C is equivalent to the conjunction of AN, NMD, and D
- Agent Neutrality
 - Nozick argues that the consequentialist approach to incorporating rights fail to capture what is crucial about the moral significance of rights, namely that rights act as 'side constraints' against our pursuit of the good
 - So you cannot pursue the violation of one even if it means fewer violations overall $\rightarrow$ i.e. killing Voldemort to prevent him from killing others
 - Nozick's theory is one that is agent-relative
 - A theory is agent-relative iff it gives different aims to different agents
 - The side constraint rights theory gives me the aim of seeing to it that I don't violate any rights, but gives you the aim of seeing to it that you don't violate any rights
 - Some moral theories say that agents have a special obligation to promote the well-being of their own family and friends, over and above the obligation they have to promote the well-being of people generally
 - Proof that if R represents an agent relative theory, then it violates AN
 - Suppose R represents the side constraint rights theory. And let worlds w_m and w_y be such that in w_m , you violate some right and I violate none, but in w_y I violate some right and you violate none. Since the theory says I ought to choose w_m over w_y , and you choose w_y over w_m , we have $R_{\{me\}}(\{w_m, w_y\}) = \{w_m\}$ and

$R_{\{you\}}(\{w_m, w_y\}) = \{w_y\}$. And this violates AN, because by hypothesis, $w_m \neq w_y$

- Reply: devise agent-relative theories of the good – what's good relative to one person need not be good relative to another
 - But such a strategy is incompatible with the idea of maximising the good – there is not a single thing which every agent ought to maximise; rather there are many things, the good relative to me, the good relative to you etc
 - What differs from person to person is not what the person ought to maximise it but how the person maximises it → we all ought to maximise the same thing, the good relative to the agent, but we each do so in our own way
 - Why can't we construct an indicator function then
- There might be an indexical content way of representing → both want to win the race but winning race depends on person i or j crossing
 - Some hold the view that the proposition "I have a beard" express a common content of centred worlds → content is a set of centred worlds: the set of every world-individual pair $\langle w, i \rangle$ such that i has a beard in w
 - Create a world order $\geq$ such that $\langle x, i \rangle \geq \langle y, j \rangle$ iff the number of rights violated by i in x is no less than the number violated by j in y
 - Then, we can define $\text{grt}_{\geq, i}(X) = \{x: x \in X \text{ and for all } y, \text{ if } y \in X \text{ then } \langle y, i \rangle \geq \langle x, i \rangle\}$
- But this way then leads to R maximises many different complete world orders → the consequentialism becomes trivially false and self-contradictory
 - If $\geq'$ is the same as $\geq$ except for the relative positions of these centred worlds, then R maximises $\geq$ iff it maximises $\geq'$. So R cannot maximise only one world order
 - But this is only because it is an objective consequentialism rather than a subjective consequentialism
- No moral dilemmas: a situation in which a person cannot avoid acting wrongly
 - Absolutist theories: a theory which holds that violations of rights are absolutely morally forbidden → always wrong → but our rightness function must map to a non-empty set
 - Some replies
 - Going incomplete: If $\geq$ is incomplete, then there is at least one set of worlds A with no greatest elements relative to $\geq$; that is, $\text{grt}_{\geq}(A) = \emptyset$. Thus a rightness function which violates NMD, though it cannot maximise any complete order, might maximise an incomplete order

- Then, need to define in terms of maximal elements rather than greatest elements
 - $\max_{\geq}(X) = \{x: x \in X \text{ and for all } y, \text{ if } y \in X, \text{ and } x \geq y \text{ then } y \geq x\} \rightarrow$ as great as every element that is at least as great as it is (basically, no one on top of you)
 - If incomplete, then $\text{grt}_{\geq}(X) \subseteq \max_{\geq}(X) \rightarrow$ every greatest element must be maximal but the converse need not hold
- We should define maximisation in terms of maximal elements to accommodate incompleteness: If an F is not one of the most G Fs, then some F is more than it
 - And this implies that in any choice situation, any alternative that is not worse than any other alternative is one of the best alternatives; but an alternative that is not less great than any other, relative to an incomplete order, might fail to be a greatest alternative relative to that order (though it must be maximal)
 - Thus, if maximisation were defined in terms of greatest elements, consequentialism would allow that an alternative might be wrong even though it is one of the best and that is something consequentialism cannot allow
 - Consequentialism must say that if an alternative is one of the best, then it is permissible
 - But if we allow incompleteness and define maximisation in terms of maximal elements, then consequentialism would still imply NMD but violate D
- Infinitely many alternatives: but if A is infinite, then even $\max_{\geq}(A)$ may be empty
- Dominance: suppose that in a given choice situation, words x and y are among the alternatives. And suppose in this situation, x is right and y is wrong. Then x dominates y in the following sense: y cannot be right in any situation where x is an alternative; the presence of x is always sufficient to make y wrong
 - Three ways to violate D
 - Satisficing: sometimes okay to do what is worse than the best provided it is not too much worse
 - Self and other: 'self-other utilitarianism' $\rightarrow$ allow people to sacrifice their own well-being in certain situations where this wouldn't be allowed by ordinary utilitarianism
 - An alternative x is right iff either (1) the total well-being of all people is at least as great in x as in every alternative, or (2) the total

well-being of all people excluding the agent, i , is at least as great in x as in every alternative

- Pareto: interpersonal comparisons of well-beings are meaningless (ordinality of utility functions) → pareto superiority of one alternative
 - Pareto-utilitarianism: (1) an alternative is wrong if it is pareto inferior to some other alternative, (2) otherwise it is right
 - There can be some alternative that is just better overall (Hicks-Kaldor compensatory sense)