

FHS Knowledge and Reality: Contextualism

Readings

Core Readings (required):

Keith DeRose (1999) DeRose, Keith and Warfield, Ted (1999). *Skepticism*. Oxford: Oxford University Press. Introduction David Lewis (1996) 'Elusive Knowledge' in *Australasian Journal of Philosophy* 74(4), pp. 549–567. (also appears in the DeRose and Warfield reader)

DeRose, Keith (1995). Solving the skeptical problem. *Philosophical Review* 104(1), 1-52.

Stanley, Jason (2005). *Knowledge and practical interests*, esp. chs. 5-6. Oxford: Oxford University Press.

Hawthorne, John (2003). *Knowledge and lotteries*, ch. 4. Oxford: Oxford University Press.

Further Readings:

Blome-Tillmann, Michael (2009). Knowledge and presuppositions. *Mind* 118(470), 241-29

Hawthorne, John (2003). *Knowledge and lotteries*, ch. 2 Oxford: Oxford University Press.

Nagel, Jennifer (2010). Knowledge ascriptions and the psychological consequences of thinking about error. *Philosophical Quarterly* 60, 286-306.

Nagel, Jennifer (2008). Knowledge ascriptions and the psychological consequences of changing stakes. *Australasian Journal of Philosophy* 86, 279-294.

Williamson, Timothy (2005). Contextualism, subject-sensitive invariantism and knowledge of knowledge. *The Philosophical Quarterly*, 55(219), 213-35. Ernest Sosa (2000) 'Skepticism and Contextualism' in *Philosophical Issues* 10, pp. 1–18. Robin McKenna (2015) 'Recent Work: Contextualism in Epistemology' in *Analysis* 75(3), pp. 489–503

Past Year Paper Questions

If you had been in the JCR you would know that you had hands, but since you are currently in an epistemology you know no such thing [2023-2024]

Can epistemic contextualists adequately explain why it seems plausible that we can't rule out sceptical scenarios? [2022-2023]

Is it easier to raise the epistemic standards governing assertion than it is to lower them? [2021-2022]

'If the word "know" picks out different relations governed by different epistemic standards in different contexts, then it doesn't matter whether S can truly be described as "knowing" that p in our context; all that matters is which of those epistemic standards S actually met with respect to p.' Discuss. [2021-2022]

If the epistemic standards governing assertion are lower in an informal conversation than in a courtroom or at a scientific conference, does this support epistemic contextualism? [2020-2021]

Since we still need an explanation for how we can "know" by ordinary standards, epistemic contextualism is no help in responding to scepticism.' Is this a fair complaint? [2019-2020]

Part I: The Problem of Scepticism and traditional Responses

The Standard Problem

Argument by Skeptical Hypothesis

- Argument:
 - I don't know that not-H (1)
 - If I don't know that not-H, then I don't know that O (2)
 - Therefore, I don't know that O. (3)
- Where H is the skeptical hypothesis, and O is some ordinary proposition like I have hands
- Some common skeptical hypotheses
 - Dreaming: Descartes's hypothesis in meditation 1 that we are in an extremely vivid dream
 - Evil Demon: that there is an evil demon who is all-powerful and feeding us with illusions about the external world
 - Russell's hypothesis: Bertrand Russell argued that the world could've come into existence mere minutes ago, but the world was complete with all the evidence of great age (i.e. our false "memories")
 - Brain-in-a-vat hypothesis: similar to evil demon but more modern

The Moorean response

Moorean facts

- Moorean facts are facts/conclusions that we are thought to know better than premises or thought to be definitely true or intuitions that we refuse to give up
 - For instance, the fact that we have hands and feet
- Lewis, "a Moorean fact...It is one of those things that we know better than we know the premises of any philosophical argument to the contrary"

Moore's argument

- Moore considers the skeptical argument of Russell's to the conclusion that he does not know "that this is a pencil or that you are conscious"
 - Moore asks "Is it, in fact, as certain that all four of these [premises for the skeptical argument] are true, as that I do know that this is a pencil or that you are conscious? It seems to me more certain that I do know that this is a pencil or that you are conscious, than that any single one of these 4 [premises] are true, let alone all four."
 - Moore accepts the intuition that these premises supporting the sceptical argument are true, but he argues that the truth of these premises are not as irrefutable as the fact that "that is a pencil"
- Moorean shift: because he takes the conclusion (3) to be a Moorean fact, he chooses to use that as a premise in his argument
 - If I don't know that not-H, then I don't know that O
 - I know that O

- Therefore, I do know that not-H
- Moore's argument buttresses on the fact that out of the set of assumptions we make, we can deny the most implausible of them all, which for him is the fact that we do not know that we are conscious or have hands
 - To further bolster Moore's argument requires us to find either (1) positive support for our intuitions or (2) damage control by providing deflationary explanations for why we have the misleading intuition we do have about the plausible statement that one chooses to deny
 - This forces us to enter the messy arena of intuitions

Worries against Moore

- Purely an intuition game so not very satisfying
- Very difficult to find arguments for comparative favourability of one fact over another since they are based purely on intuition and seem true to most of us
- Moore overlooks the promising possibility that one can hope for an explanation of how we fell into the puzzle of conflicting intuitions in the first place and we can explain how premises that together imply a conclusion that we find so incredible can themselves seem so plausible to us
 - Such an explanation can take the form of explaining, for the member of the set that one seeks to deny, why it seems to us to be true, though it's in fact false
 - The game would then be a game of damage control – of providing a deflationary explanation for why we have the misleading intuition we do have about the plausible statement one chooses to deny → if this is the case, then it might seem that the skeptical argument does best

Putnam's Argument from Semantic Externalism

Semantic Externalism

- The thesis that the contents of one's thoughts are not completely determined by internal fact about what's going on inside one's head, but are at least partially determined by such "external facts" such as the nature of the items one has been in contact with
 - For suppose an ant coincidentally draws an outline of William Churchill. We would not say he is drawing a picture of William Churchill
- Putnam denies what he calls "magical theories of reference"
 - According to such a theory, some representations such as names are believed to have a magical connection with their bearers
 - But once we realise that this connection is achieved only via context and contingency and convention, the mystical nature of this connection disappears

- Mental representations face the same problem as physical representations → they no more have a necessary connection with what they represent than physical representations do
- Suppose humans have never seen a tree and they have never imagined trees. Imagine that one day a picture of a tree is accidentally dropped on their planet by aliens.
 - For us the picture is a representation of a tree
 - For these humans, the picture only represents a strange object
- Suppose then one human has a mental image which is exactly like one of my mental images of a tree as a result of having seen the picture
 - His mental image is not a representation of a tree but a representation of the strange object that the mysterious picture represents
 - Imagine that the picture was not caused by a real tree
- Same thing is true of words → if a discourse on paper might seem to be a perfect description of trees, but if it were produced by monkeys randomly hitting keys on a typewriter for millions of years, then the words do not refer to anything
- Putnam's example of trees: in particular, you cannot think about trees if you haven't been causally connected with trees in the proper way

Putnam's argument

- From semantic externalism, Putnam says that because you haven't been causally connected with trees in the proper way, then you cannot be properly said to be thinking about trees
- Analogously, if we accept semantic externalism, when such a BIV thinks such thoughts as those she expresses via the sentences, "There's a tree", or "Here's a hand" or "I am not a BIV", then it is not thinking the same thing that those words would express in our mouths/minds (given that we're not BIVs), since the BIV is not causally connected with hands, trees, vats etc in the way needed to have such thoughts
 - These words, said by the BIV, refers instead to imitations of the trees in the image or the electronic impulses that causes tree experience
- The semantic externalist only need to be committed to the negative thesis that the word "trees" do not refer to real trees
- Putnam concludes that if we are BIVs, then "what we now mean by 'we are not brains in a vat' is that we are not brains in a vat in the image or something of the kind (if we mean anything at all)
 - So, if the BIVs succeed in meaning anything true or false by their relevant thoughts involving "tree", "hand", "vat", and so on, they mean one of the

proposals of electronic impulses or image, or some closely related other thing on which the relevant thoughts turn out to be true

- First strategy against SA
 - If we are BIVs, then by “I am not a BIV”, one means that one is not a BIV-in-the-image is obviously true.
 - If we are not BIVs, then by “I am not a BIV”, one means that one is not a BIV, which in that case is true.
 - So either way it is true
- Objections against first strategy
 - Seems to only yield the conclusion that my utterance of “I am not a BIV” is true when we need I am not a BIV to be true
 - It also requires the strong, positive externalist thesis that by her use of such sentences as “I am not a BIV”, a BIV means that some true thought such as that she is not a BIV-in-the-image, and it is unclear whether the thought experiments used to support semantic externalism support such a conclusion
 - Putnam only claimed the weaker, negative results about what a BIV could not mean
- Second strategy against SA
 - The externalist thesis used on this second strategy is only a negative claim about what a BIV cannot mean or think – that by “tree”, “hand”, “vat” and so on, the BIV does not refer to the real trees, hands, vats etc due to her lack of causal connection
 - This negative thesis is combined with a positive claim to the effect that we do have the thoughts in question – the thoughts that the BIVs cannot have → these together imply that we are not BIVs
- Objections against second strategy
 - Does the antiskeptical have any right to the second claim? Seems just a more sophisticated form of a Moorean response

Worries against Putnam (General Objection)

- Against semantic externalism: what about numbers or imaginary objects? We have no causal connection with them
- Recently evaded cases: if Putnam’s argument is robust, it is only robust against certain SAs and not all → for instance, suppose that we have only been recently evaded and thus actually we do have causal connections with trees, hands etc
 - I do actually mean tree, vat, hand when I express those → falsely thinking “I am not a BIV”
- Putnam concedes too much to the skeptic: by proving that one is not a BIV, one seeks, it seems, in following this strategy, to make it the case that the first premise of the Argument by Skeptical Hypothesis is false as applied to oneself

- This argument seems to be “heroic” in that only the antiskeptical and others that follow him who know and understand the argument are helped by the argument
- People who have never encountered Putnam’s proof may be said to not know that they have hands because they do not know that the premise “I do not know that not-H” is actually false when they express it

Nozick’s Denial of Closure

The Principle of Closure

- We can infer from “S knows that P” and “S knows that P implies Q” that “S knows that Q”
- The Principle of Closure is said to underwrite almost all Skeptical Arguments → in fact, its premise 2 is said to rest on it
 - Proof: By assumption, I know that O implies not-H. By closure, if I know that O, then I know that not-H. By contraposition, if I do not know that not-H, then I do not know that O
- So rejecting closure amounts to rejecting premise 2

Nozick’s subjunctive conditionals and denial of closure

- Nozick’s account of knowledge includes the standard truth and belief conditions as well as the following subjunctive conditionals
 - If P weren’t true, S wouldn’t believe that P
 - If P were true, S would believe that P
- For Nozick, the first subjunctive helps us against the skeptical hypothesis → the belief that you’re not a BIV fails to meet the condition that if you’re belief weren’t true (i.e. that you were a BIV), you would still believe that you weren’t a BIV
 - The BIV has all the sensory experience and everything that you believe including that you aren’t a BIV
 - But the belief of ordinary propositions such as you have hands does satisfy these conditions because if you did not have hands, then you would not believe that P
- So Nozick’s subjunctive conditionals amounts to denying closure: there is a case in which you do not know that not-H (i.e. you do not know that you are not a BIV), while you know that O (i.e. you have hands)
 - The closest possible world in which you do not have hands is not one in which you are both a BIV and do not have hands by construction, but the world in which you simply do not have hands
 - But surely this closeness relation depends on whether you are in a BIV in the first place or not → inevitable circularity

Worries

- The main worry is the extremely strong intuitions we have for the principle of closure
 - Both Drestke and Nozick admit that the principle of closure is extremely intuitive and Nozick compares it to a steamroller

Part II: Contextualism

Summary of Contextualism by DeRose

Two senses of knowing

- The immediate reaction to the skeptical hypothesis argument is that only in philosophy do we have such demanding standards of knowing
 - There is a disjunct between ordinary, non philosophical uses of the word “know” and philosophical uses of the word know
- The main goal is to not make such a distinction ad hoc between philosophical and non philosophical contexts and to find actual non-philosophical uses of the word knows that has a lower standard

The Contextualist Strategy

- Current contextualists posit a wide variety of epistemic standards, based on rules by which what is said (Lewis) in a conversation that can change the standards that are in place
- They concede that when the skeptic presents their argument, that through one of the rules exploited, they raise the standard for knowledge such that their argument actually follows, thus making our conclusion that we “don’t know” true
 - Benefit is that they get to uphold closure, relative to any set standard for knowledge
- The main thrust is that the skeptic employs higher standards for knowledge than usual
- Two main theorists for contextualism: Lewis and DeRose
 - Lewis: focuses more on eliminating relevant alternatives → S knows P iff S’s evidence eliminates every possibility in which not-P, except for those possibilities that are properly ignoring
 - Must determine what are irrelevant alternatives; Lewis thinks that relevant alternatives that we need to eliminate as knowers are those in which not-P that we’re either paying attention to or else are improperly ignoring
 - Rules-based approach: the Skeptic, on Lewis’s account, exploits “the Rule of Attention”, according to which those alternatives we’re paying close attention are relevant (i.e. the BIV hypothesis)
 - By calling attention to various possibilities that we typically and properly ignore, the skeptic makes those alternatives relevant and if our evidence fails to eliminate them (and they do), then we’re vulnerable to the Skeptics’ higher standard for knowledge
 - DeRose: focuses more on what he finds correct in Nozick while avoiding its pitfalls → claiming that the conversational rule that the skeptic exploits in SA is the “Rule of Sensitivity”
 - Where a belief is sensitive is if it meets Nozick’s first subjunctive condition → that if not-P, then you would not have believed that P

- Rule of Sensitivity: When it is asserted that some subject S knows (or does not know) some proposition P, the standards for knowledge tend to be raised, if need be, to such a level as to require S's belief in that particular P to be sensitive for it to count as knowledge
- Since, as DeRose argues, you must be in a very strong epistemic position before any belief you might have that a skeptical hypothesis is false can be sensitive, the skeptic's assertion "I don't know that not-H", will by the rule of sensitivity, raise the standards for knowledge to a level at which you could as knowing either that not-H nor that O
- The rule of sensitivity explains why we won't call someone a knower if we think his/her belief is not sensitive, and thus explains why Nozick's account of knowledge can seem so attractive
- But accepting this does not violate the principle of closure that Nozick seems to be committed to
- General issues
 - Concedes too much to the skeptic since these strategies allow that we don't know according to the high standards of SA
 - But I don't really think this is a problem

DeRose's Contextualism

Skeptical Arguments

- Argument from Ignorance
 - I do not know that not-H
 - If I do not know that not-H, then I do not know that O
 - Therefore, I do not know that O
- Argument from Possibility
 - It is possible that H
 - If it is possible that H, then it is possible that not-O
 - So, it is possible that not-O
 - If it is possible that not-O, then I don't know that O
 - So, I don't know that O
- The strategy of rejecting the skeptical argument should (i) yield an explanation of how we fell into this skeptical trap in the first place, and not settle for making a simple choice of rejecting one of the statements of SA, while (ii) allowing us to have substantial knowledge about ordinary propositions

Basic strategy of Contextualism

- According to contextualist theories of knowledge attributions, how strong an epistemic position S must be in with respect to P for A's assertion to be true can vary according to features of A's conversational context

- The skeptic manipulates the semantic standards for knowledge thereby creating a context in which she can truthfully say that we know nothing or very little
- DeRose assumes a skeptic-friendly version of contextualism – one according to which the philosophical skeptic can (fairly easily), and does succeed in raising the standards for knowledge in such a way as to make her denials of knowledge true
 - Some contextualists may think that it's not so easy to so raise the standards for knowledge and that a determined opponent of the skeptic can keep the standard low
 - Important point is to identify the mechanism by which the skeptic at least threatens to raise the standards for knowledge
- The contextualist's ultimate point will then be: to the extent that the skeptic does succeed, she does so only by raising the standards for knowledge, and so the success of her argument has no tendency to show that our ordinary claims to know are in any way defective
- DeRose highlights that the contextualist's solution, even though communicated in terms of conversational contexts, will have to be generalised to explain why the argument can be so appealing even when one is considering it in solitude, with nothing being said
 - First form of generalisation: it can be maintained that there is a rule for changing the standards for knowledge that governs the truth conditions of our thoughts regarding what is and is not known that mirrors the rule for the truth conditions of what is said regarding knowledge
 - Second form of generalisation: our judgement regarding whether something can or cannot be truly asserted might be held to affect our judgement regarding whether it's true or false, even when we make this judgement in solitude, with nothing being said at all
- DeRose thinks that contextualists have yet to explain the persuasiveness of AI and thus have yet to solve the puzzle, if we haven't located and explained the conversational rule of mechanism by which the skeptic raises

DeRose's Contextualism: Rule of Sensitivity

- DeRose draws on Nozick's theory of needing sensitive belief to explain why SA is so compelling and why we don't know that the hypothesis is ever false – DeRose calls this the subjunctive conditionals account (SCA) of the plausibility of SA's first premise
 - According to SCA, the problem with my belief that I'm not a BIV – and I do have such a belief, as do most of us – is that I would have this belief even if it were false → this makes it hard to claim to know that I'm not a BIV
 - S's belief that P is insensitive if S would believe that P if P were false
- Also useful in explaining our reluctance to ascribe knowledge in certain lottery situations even where the odds of your being a loser are astronomically high

- Your belief that you're a loser is one that you would hold even if you were the winner
- Problems with SCA: there are cases in which it seems to us that S does know that P even though we judge that S would believe that P even if P were false
 - But DeRose says that this general inclination need not be exceptionless to perform the explanatory role SCA assigns in it
 - Problem Case: A grandmother sees her grandson is well when he comes to visit; but if he were sick or dead, others would tell her he was well to spare her upset. Yet this does not mean she doesn't know he is well when she sees him
 - Here it seems that the grandmother knows her grandson is well even though even if he was in fact sick, she would still believe that he is well
 - Nozick attempts to relativise the sensitivity condition to method of knowing: if p weren't true and S were to use M to arrive at a belief whether that p, then S wouldn't believe, via M, that p
 - But DeRose suggests that context takes account this relativising of method: even to one aware of the likelihood that the grandmother's family would have kept her in the dark about her grandson's condition were he not well, it can seem like Nozick's initial condition is met → it seems that if the grandson were not well, the grandmother would not believe he was well, since she is looking right at him
 - The reason for this is that when one searches for the possible worlds most similar to the actual world in which the grandson is not well, the respects in which the possible worlds are to resemble the actual world is a highly context-sensitive matter – especially where the context focuses on one's attention on the grandmother and her cognitive and recognitional abilities, one can place heavy weight upon similarity with respect to the method she is using to arrive at her belief
- DeRose argues that SCA is confirmed once we figure out what makes Skeptical Hypotheses not convincing
 - Hypothesis: the zoologists tells you that you can't know that that is a zebra because you wouldn't know if it is a painted mule or not
 - Investigations that would seemingly allow us to know that animals aren't painted mules would also affect our judgments as to the truth value of the subjunctive conditional so critical to SCA
 - So your belief can still be sensitive
- But DeRose rejects Nozick's denial of Closure which results in the abominable conclusion that "S knows that O" even though "S does not know that not-H"
 - This is simply too abominable to accept → DeRose proposes a very strong endorsement of the closure principle underwriting the second premise of SA: "If I do not know that not-H, then I do not know that O"

- The reason for this is that we are in at least as good a position to know that the hypothesis is false as we're in to know the targeted piece of presumed ordinary knowledge
 - So if I do not know that I am not in a BIV, then my epistemic position is also not strong enough to know that I have hands
- The second premise leads us to the following insight regarding skeptical hypotheses: One's epistemic position with respect to propositions to the effect that skeptical hypotheses don't hold must be stronger than it is with respect to other, more ordinary propositions if belief in such propositions is to be sensitive
- Possible Worlds:
 - An important component of being in a strong epistemic position with respect to P is to have one's belief as to whether P is true match the fact of the matter as to whether P is true, not only in the actual world, but also at the worlds sufficiently close to the actual world → one's belief should not only be true, but also non-accidentally true where this requires one's belief as to whether P is true to match the fact of the matter at nearby worlds
 - The further away one can get from the actual world while still having it be the case that one's belief matches the fact at worlds that far away and closer, the stronger a position one is in with respect to P
 - If the truth-tracking of one's belief as to whether P extends far enough from actuality to reach the closest not-P worlds, then one doesn't believe that P in those closest not-P worlds, and one's belief that P is sensitive
- Example: Consider my belief that I have hands. I believe this at the actual world and it's true. In other nearby worlds in which I have hands, I believe that I do. There are also alarmingly close worlds in which I do not have hands, perhaps because of a car accident. In the closest of these non-P worlds, I am now fully aware of the fact that I am handless and my belief as to whether I have hands matches the fact of the matter
 - The case in which I am a BIV is very distant → while there are closer worlds in which the match fails, it seems that in a fairly wide range of worlds surrounding the actual world, my belief as to whether I have hands does a good job of matching the fact of the matter
 - Thus, I am in a strong epistemic position with respect to the matter
- Example: Now let my belief be I am not a BIV. Where not-P is quite remote, one can be in quite a strong epistemic position with respect to P merely by believing that P in all nearby worlds. As I do believe this P in such nearby worlds, I am in a strong epistemic position with respect to this P. This can occur, and in my case, does occur, even though one's belief as to whether P doesn't match the fact of the matter in the closest not-P worlds: Since even the closest of the not-P worlds are quite distant, one's belief as to whether P needn't match the fact of the matter that far from the actual world for one to be in a quite strong position with respect to P

- But for one's belief that P to be sensitive, one must not believe that P in the closest not-P worlds. Since Skeptical Hypotheses tend to fasten on somewhat remote possibilities, then one can be in a relatively strong position with respect to beliefs to the effect that they don't obtain while these beliefs remain insensitive
- By contrast, where P is such that there are both P and not-P worlds very close to the actual world, one's belief that P must be sensitive in order for one to be in even a minimally strong epistemic position with respect to P, and conversely, one needn't be in a very strong position for one's belief to be sensitive
- Fundamental point: one's epistemic position with respect to propositions to the effect that Skeptical Hypotheses don't hold must be stronger than it is with respect to other propositions before beliefs in such propositions can be sensitive
 - Suggests that new contextualist account of how skeptic raises the standards for knowledge: when it is asserted that some subject S knows (or does not know) some proposition P, the standards for knowledge (the standards for how good an epistemic position one must be in to count as knowing) tend to be raised, if need be, to such a level as to require S's belief in that particular P to be sensitive for it to count as knowledge
 - Where the P involved is a skeptical hypothesis, then this rule dictates that the standards will be raised to a quite high level, for as we've seen, one must be in a stronger epistemic position with respect to a proposition stating that a skeptical hypothesis is false, relative to other, more ordinary propositions, before such a belief in such a proposition can be sensitive
 - Context determines how strong an epistemic position one must be in to count as knowing → contextual determined sphere of possible worlds, centred on the actual world
- The Skeptic raises the standards of knowledge by choosing a H s.t. (i) we will be in at least as strong a position to know that not-H as we're to know that O, but (ii) any belief we might have to the effect that not-H will be an insensitive belief

Lewis's Epistemic Contextualism

Central Thesis: advances the typical contextualist argument that the truth of the ascription of "S knows that p" depends on the features of the conversational contexts; Lewis, in particular, advances a rules based account

Lewis's modal epistemology

- Fallible knowledge i.e. knowledge despite uneliminated possibilities of error sounds contradictory
 - But we have very little 'traditional' infallible knowledge
- Knowledge must be infallible yet we have fallible knowledge or no knowledge at all

- But Lewis says that fallibility is better than skepticism
- We seem to know a lot in daily life but not when we're analysing the theory of knowledge i.e. doing epistemology
 - When I am in the supermarket, I seem to know that I have hands, but once I am in an Oxford tutorial, I do not know that I have hands
- Context-dependency is not an exotic thesis: "it's evening" is truly said only when it is said in the evening → different contexts evoke different standards of evaluation
 - "Essendon played rottenly" vs "Easybeats played brilliantly" and yet Essendon won → because "rottenly" and "brilliantly" are subject to the contexts of Essendon (perhaps below their standard) and Easybeats (perhaps above their standard)
 - Or "Magnus Carlsen played rottenly while I played brilliantly, even though I lost decisively"
- Lewis thinks that no justification is good enough
 - Your true opinion that you will lose the lottery isn't knowledge, whatever the odds
 - The greater the number of losing tickets, the better is your justification for believing you will lose; yet there is no number great enough to transform your fallible opinion into knowledge – since you might just win
 - No justification is good enough short of deductive argument
- Justification is also unnecessary: what non-circular argument supports our reliance on perception, on memory, on testimony?
 - Assumes internalist justification
- Lewis proposes the following for knowledge
 - S knows P iff P holds in every possibility left eliminated by S's evidence; equivalently iff S's evidence eliminates every possibility in which not-P
 - S knows P iff S's evidence eliminates every possibility in which not-P, except for those possibilities that we are properly ignoring
- The question is what counts as possibilities?
 - Cannot limit ourselves to real possibilities that conform to actual laws of nature and also to actual past history
 - Propositions about laws and history are contingent and may or may not be known
 - Cannot limit ourselves to 'epistemic' possibilities for S – possibilities that S does not know to not obtain as it would drain the definition of content
 - Assume that knowledge is closed under strict implication. Then knowledge of a conjunction is equivalent to knowledge of every conjunct. P is the conjunct of all propositions not-W, where W is a possibility in which not-P. This yields the equivalence: S knows that P iff, for every possibility W in which not-P, S knows that not-W iff every

possibility which S does not know not to obtain is one in which P iff P holds throughout S's epistemic possibilities

- To turn this into a substantive definition, we need to say that S's epistemic possibilities are just those possibilities that are uneliminated by S's evidence

What counts as properly ignored uneliminated alternative possibilities?

- Rule of Actuality: the possibility that actually obtains is never properly ignored
 - Nothing false may be presupposed
 - It follows that only what is true is known, hence Lewis does not include truth as part of the definition of knowledge
 - Whose actuality? The actuality of the subject of our ascriptions is different from our actuality → When I ask what Fred knew yesterday, even though Fred and I are worldmates, when we build subject and time into the possibilities *de se et nunc*, then Fred's actuality is the possibility *de se et nunc* of being Fred on Sept 19 and my actuality is the possibility *de se et nunc* of being Matthew on Sept 20
 - Sometimes we also ascribe knowledge to others in possible worlds
 - Basically, ascriber's actuality
- Rule of Belief: a possibility that the subject believes to obtain is not properly ignored, whether or not he is right to so believe
 - This is not what Lewis thinks the skeptic uses since the skeptic doesn't necessarily believe himself to be a BIV; only that he does not know
 - Degree of belief must be sufficiently high → Lewis thinks this depends on how much is at stake; when error would be disastrous, few possibilities may be properly ignored
 - But even at high stakes, some possibilities may still be properly ignored: jurors can ignore the possibility that it was the dog who shot a murder victim even though wrong conviction is disastrous
- Rule of Resemblance: Suppose one possibility saliently resembles another, then if one of them may not be properly ignored, neither may the other
 - Explains why you do not know that you will lose the lottery no matter the odds, since the possibilities in which you lose and win the lottery are saliently similar to one another → either ignore both or neither
 - And since one of them actually obtains, you can't ignore that one, and so you must not ignore either
 - Also solves the Gettier problem
- Rule of Reliability: we can presuppose what usually work to work – a possibility in which reliable processes fail may be properly ignored
 - Visual experience etc
- Rules of method: we are entitled to presuppose – very defeasibly – that a sample is representative; and that the best explanation of our evidence is the true explanation

- Entitled to possibly ignore failures in those two standard methods of non-deductive inference
- The general rule consists of a standard dispositions to presuppose reliability in whatever particular case may come before us
- Rules of Conservatism: Suppose that those around us normally do ignore certain possibilities, and it is common knowledge that they do, then very defeasibly, these generally ignored possibilities may properly be ignored
- Rule of Attention: when we say that a possibility is properly ignored, we mean exactly that and not that it could have been properly ignored
 - A possibility not ignored at all is ipso facto not properly ignored i.e. the mere fact of giving a possibility attention makes it not possible for us properly ignore it
- Rules can defeat one another
 - Rule of Reliability may be defeated by Rule of Actuality

Sosa's Objection to Contextualism as a Strategy against Skepticism

Central Thesis: Contextualism does not seem to be an epistemological, but instead a linguistic thesis and thus may not be applicable to the problem of skepticism

- Central question: supposing epistemic vocabulary to be correctly applicable in contexts that set a different threshold from that of epistemological inquiry, how relevant can that be to epistemological questions about the nature, conditions, and extent of human knowledge
- Contextualism usually features the following thesis: Two people affirming the same such sentence may yet be right and wrong respectively, owing to different contexts of use
 - The contextualists concedes that in a context of philosophical reflection it is false to say "I know I have a hand" while insisting that in ordinary contexts, in the home or the marketplace or the sports arena, it is not false but true to say that same thing
- Contention: even if the utterance of such a sentence is often enough true, what bearing might that have on epistemology
 - The fact that People often utter truths when they say 'Someone loves me' does not have much relevance to the worry "Does anybody at all love me?"
 - Analogously, the fact that people often utter truths when they say "I know there are hands" does little to bear on the question "Do people ever know that there are hands"
 - Just because people utter truths about the fact that "they know they have hands" doesn't answer the deeper philosophical question about whether they have knowledge about the external world

- The truth of the knowledge claim in ordinary contexts does not always logically entail knowledge → most be so in the robust epistemic sense and not merely coincidentally
 - Does not tell us in what sense the conditions for knowledge are met
 - Furthermore, the first sentence does not even entail that people often utter truths when they say that they know there are hands → since it will not follow that people ever do claim, in an ordinary context, that they know things, as opposed to making utterances of the form “I know such and such”
- The contextualist line deriving from epistemic contextualism hence does not much support, for philosophers, the claim that people do in ordinary contexts after all know things nor does it even much support the claim that speakers are often enough right when they say that people know things
 - Limits the epistemological interest and relevance of EC-contextualism
- Objection: Contextualists do not just propose EC but specify what contextual factors affect the setting of the relevant thresholds, yielding results as to how the vocabulary of knowledge is correctly applicable in our context of philosophical inquiry
 - Sosa only says that EC is limited in specified ways → it does not follow from the contextualist position, we often enough know ordinary propositions in ordinary contexts → only our utterance “I know that p” is true
 - It does not bear on the question whether I’ve specifically met the deep epistemological criterion of knowledge
 - On occasion, the term “know” is made so easily satisfiable that in some ordinary contexts, if someone is very sure that p, it makes it true to say they “know” that p → and this is clearly not relevant to our concern to understand the nature, conditions, and extent of knowledge that we have

McKenna’s Summary of Arguments for and Against Contextualism

Contextualism

- Primarily focused on knowledge ascriptions i.e. utterances of sentences like “John knows that he has hands”
- Context that matters is the conversational context, not the social or historical context
- The context that matters is the ascriber’s, not the subject of the ascription

Fundamental Questions about Contextualism

- What does it mean to say that the truth-conditions of S’s knowledge ascription depend on the epistemic standard operative in her context?
 - Contextualists think that sentences containing the expression ‘knows’ can express different propositions in different contexts of utterance
 - Imagine S1 being in context E1, S2 being in context E2, and S1 and S2 both utter “S* knows that p”, the proposition expressed by S1 is something like “S*’s epistemic position with respect to p is strong enough to E1” whereas the proposition expressed by S2 is something like “S*’s epistemic position with

respect to p is not strong enough to satisfy E2”

- Is the epistemic standard operative in her context just her standard, or some other standard?
 - Most contextualists think that the operative standard need not be the ascriber’s standard (like DeRose)
 - The operative standard depends on the conversational context whether because it’s obvious the speaker and interlocutor all accept them, or because they have been explicitly accepted
- What are the alternatives to contextualism
 - Relativism: the thesis that there are multiple contexts to evaluate one’s knowledge ascription → a single proposition is true relative to some contexts of assessment and yet false relative to others
 - Contextualism fixes the assessment context to the ascriber’s
 - Invariantism: the sentence of the form “S knows that p” expresses the same proposition in every context → the truth depends on sorts of things epistemologists have always talked about – evidence, reliability of perceptual and other cognitive faculties
 - Strict invariantism has two components
 - Invariantist component: knowledge ascriptions express the same proposition in each context
 - Strict component: whether someone ‘knows’ just depends on truth conducive factors

Arguments for Contextualism

- Contextualists appeal to a putative general phenomenon: the appropriateness of ascribing ‘knowledge’ depends on the context of ascription
 - Low: Hannah and her wife Sara are driving home on a Friday afternoon and plan to stop at the bank on the way home to deposit a cheque. It is not important that they do so since they have no impending bill and Hannah says “I know that the bank will be open tomorrow. I was just there two weeks ago on Saturday morning. So we can deposit our cheque tomorrow morning.”
 - High: Same as Low but they have an impending bill so it is very important that they deposit their cheque by Saturday. Hannah notes that she was at the bank two weeks before on a Saturday morning, and it was open but as Sartah points out banks do change their hours, so Hannah says, “I guess you’re right. I don’t know that the bank will be open tomorrow.”
 - This argument is highly criticised
 - Not all experimental studies agree on the relevant factors → practical factors sometimes feature more strongly
 - Other views can explain the general phenomenon → Grice distinguishes between appropriate and true → it may be true that X while not being appropriate to say that X

- Skepticism
 - Contextualists can explain why we find the skeptical argument so compelling in the first place and allows us to retain ordinary knowledge in ordinary contexts without violating any intuitive principle
 - Worries
 - Concedes too much to the skeptic, especially for Lewis → the skeptic is one who pushes for these standards and Lewis seems to allow the skeptic inevitably succeed → as soon as the argument is run, she triumphs
 - Contextualism is irrelevant to the skeptical problem: the contextualist says that, while we don't 'know' according to the high standards operative in skeptical contexts, we do 'know' according to the low standards operative in more standard contexts → but this fails to engage with the sort of radical skeptic who argues that we don't 'know' according to any standards because, epistemically speaking, all beliefs, whether true or false, are on par → to the extent that skeptical arguments are intended to establish this conclusion, the problem with them has little to do with contextualism

Arguments against Contextualism

- Linguistic objections
 - Some contextualists have proposed that 'knows' is analogous to gradable adjectives like 'tall' or 'flat' (DeRose)
 - A gradable adjective like 'tall' makes implicit reference to a contextually provided scale of height → similarly, one might think that 'knows' makes implicit reference to a contextual provided epistemic standard where we can think of particular epistemic standards as points on a scale
 - First, unlike 'knows', gradable adjectives accept various degree modifiers (completely/very/quite) and comparative constructions
 - X is flatter than y
 - S knows that p more than y
 - Second, 'knows' doesn't accept clarifying devices
 - X is roughly flat
 - X roughly knows that p
- Problem from Disagreement: contextualists are unable to make sense of disagreement
 - Consider the following conversation:
 - Ailsa: I know that the animal in the cage is a zebra
 - Laurie: No, you're wrong. You don't know that it's a zebra because you can't rule out the possibility that it's a mule cleverly disguised to look

like a zebra

- Context 1: Niall and Cormac at the zoo
- Niall: Do you know what animal that is? Cormac: I know that it's a zebra
- Context 2: Niall is telling Laurie about his trip to the zoo.
- Laurie: What animals did you see?
- Niall: We saw a zebra
- Laurie: How do you know that it was a zebra?
- Niall: Cormac told me it was because it had black and white stripes.
- Laurie: Cormac's wrong, he doesn't know that. How could he rule out the possibility that it was a mule cleverly disguised to look like a zebra
- Ailsa and Cormac meet the standards in their respective contexts and they weren't wrong about anything: makes it hard to explain Laurie's and Cormac's uses of the disagreement marker 'you're wrong'
- But not clear what to take from this → adopt relativism, invariantism?
- Assertion and context
 - Williamson argues that knowledge is the norm of assertion
 - Assertion: S is in a strong enough epistemic position to assert p iff S knows p
 - DeRose appeals to the knowledge norm of assertion in arguing for contextualism with the following modification
 - Assertion*: S is in a strong enough epistemic position to assert p iff 'S knows p' is true in S's context
 - But Hawthorne argues that assertion causes problems for contextualism
 - Imagine Sarah is deciding whether to go into the bank to cash her cheque now or come back tomorrow (Saturday). There's no rush to cash the cheque, and Sarah has good evidence that the bank is open tomorrow (she remembers it being open on previous Saturdays). Sarah decides to come back tomorrow. Hannah is making the same decision – while Hannah has the same evidence as Sarah, it is imperative that Hannah cash her cheque before Monday. Hannah goes in to cash her cheque. According to the contextualist, Sarah can truly say that both of them 'know', whereas Hannah can truly say that neither of them 'knows'. But according to Assertion*, Sarah is in a strong enough position to assert while Hannah isn't. So the following are true
 - Hannah: Sarah doesn't know that the bank is open, but she's in a strong enough epistemic position to assert that the bank is open
 - Sarah: Hannah knows that the bank is open, but she isn't in a strong enough epistemic position to assert that the bank is open
 - This violates the Norm of Assertion

Hawthorne's Objections to Contextualism

Violates the following intuitive principles

- Assertion Constraint
 - Relevant normative facts do not seem to be ascriber dependent
 - From the perspective of a contextualist: Consider that I, the ascriber, am in a context in which I can truly say "You know you will not have enough money to go on an African safari this year." And unbeknownst to me, you have from this deduced and thereby come to believe that you will not win the lottery
 - You assert without me hearing: I will not win the lottery; In my context, the following ascription is correct: You know you will have not enough money to go on an African Safari this year
 - Assuming Single Premise Closure, the following is true according to my context
 - Anything that you have competently deduced from knowing that you will not have enough money to go on an African Safari this year and thereby come to believe it is something you know as well
 - But it may well be that it was outright inappropriate for you to have asserted what you did five minutes ago i.e. I will not win the lottery
 - You ought not to have asserted what you asserted is true in my mouth (whether I know or not it is true)
 - This results in: there are things people know but ought not to assert because their epistemic position is not strong enough with respect to those things
 - My standards for 'know' are higher than S's, so that S stands in the relation that he expresses by 'know' to some p that he asserts, but that he does not stand in the relation that I express by 'know' to p
- Single Premise Closure
- The Practical Reasoning Constraint
 - Suppose that I have high standards s.t. "You do not know that you are going on vacation is true in my mouth" but it would seem that you may be in a practical environment where it is perfectly appropriate (by anyone's standards) to use the proposition that you are going on vacation as a premise in your practical reasoning
 - So the following would be true in my mouth: You should rely on propositions that you don't know to be true in your practical reasoning
- The Epistemic Possibility Constraint

Part III: Against Contextualism – Varieties of Invariantism

Hawthorne's Subjective-Sensitive Invariantism

Central Claim: two mechanisms bear on the truth of knowledge claims but that focus on subject-sensitivity. The sensitive moderate invariantist claims that the extension of know depends not only on the kinds of actors traditionally adverted to accounts of knowledge but also on the kinds of factors that in the contextualist's hands make for ascriber dependence.

- A subject-sensitive salience rule
 - If John and Mary are to properly assert 'Smith knows that the flight stops in Chicago', they will have to know that Smith knows that the flight stops in Chicago; since knowledge is factive, they will then have to know that the flight stops in Chicago. But by hypothesis, they do not know that the flight stops in Chicago
 - Cohen's claim that the case provides a proof that contextualism is the only serious option to skepticism can be challenged
 - Depends on the context of the subject instead
 - Once the possibility of winning the lottery is salient to someone (salience is not mere attention), then that person knows neither that he will lose the lottery nor that he will not be able to afford an African Safari
 - Does not violate premise: suppose someone knows that he cannot afford the African Safari, and knows that if he cannot afford the Safari, he did not win the lottery. But if the counterpossibility of winning the lottery becomes salient to him sometime during the reasoning process, then knowledge of the premise becomes destroyed
 - So Single-Premise Closure does not entail that the person knows he will lose the lottery, even if the latter can be deduced from something he currently believes or knows (since it'll get destroyed)
 - Certain inferences are anxiety provoking: when reasoning, they realise or it is made salient to them the counter-possibilities that refute what they might believe in
 - Not all counterpossibilities will be salient
- Problem: after the counter possibilities are made salient, does it licence me to say 'Up to 5 minutes ago, I knew I would not have enough money to go on the African Safari'
 - If at a later time I don't know the ordinary proposition, then I cannot assert that I used to know it, since knowledge is the norm of assertion
 - To properly assert that one used to know that p, one needs to know p now → so the statement might be in fact true, but it is not proper for me to assert it
 - So, Hawthorne claims that the awkwardness arises from my assertion, rather than the fact that the statement is true
- Problem: we are willing to say "I was mistaken in thinking that I did know that" and "I never did know that" → the sensitive moderate invariantists must account for these

intuitions that suggest the knowledge relation picks out the same ordered pair regardless of the context

- But the sensitive invariantists want to say that people do know ordinary propositions, and that considerations of the skeptical kind do not render all past knowledge moot at those times
- Hawthorne's solution of projectivism: we tend to project our own lack of knowledge to others
 - Against the contextualist: the reason why we may think that "S does not know that p" after some counterpossibility has been made salient to us is because we overestimate the likelihood of an event occurring due to availability
 - Solving the sensitive invariantist problem: we overproject our own lack of knowledge onto our past selves; it is correct to say of ourselves at this time that we do not have knowledge, but it is not properly assertible to say we lack knowledge in the past
 - I do not think this works very well
- Problem: what is stopping me from knowing things is worrying is not something we find intuitive
 - We want to accept van Fraassen's appealing Reflection Principle: expected epistemic possibility equals current epistemic probability
 - Will unduly affect thinking people → the philosopher will know nothing
- What is it for a possibility of error to be salient
 - Not merely entertaining as Lewis suggests → too permissive
 - Salient counterpossibility of error is a certain kind of intellectual seeming: it is a claim of epistemic possibility tied to the epistemic subject
 - The content of intellectual seeming associated with the salient possibility of error is basically this: for all the subject knows, p!
 - There is now nothing puzzling that the salience of the possibility of error induces retraction of a knowledge claim: just a special case of the claim that its intellectually seeming that p induces us to retract propositions incompatible with p
- Ways for counterpossibility to undermine knowledge
 - Belief Removal: destroying the belief that p destroys knowledge that p since one comes to believe not-p
 - Evidential model: I might be provided evidence for p and so if it seems to me that I might be wrong, then this provides me with evidence that I do not know p
 - Defeat: if, at t, I get evidence that I do not know that p, and I have no effective means for rebutting the evidence, then I do not know that p at t

- Objection: if someone on the train presented me with Zeno's argument that motion is impossible and while not knowing how to answer them, I choose to ignore them, do I then not know that I am moving?
 - Authority: first person authority → If I believe of myself that I do not know, then, automatically as it were, I do not know
- Epistemic possibility constraint
 - Allowing knowledge that p is compatible with a small epistemic chance that not-p
- Practical Environment Constraint
 - Salience is not enough: lottery propositions i.e. I will lose the lottery, is a bad premise for practical reasoning EVEN if counterpossibilities are not salient to me
 - I will be going to Blackpool next year; so I won't die beforehand, So I ought to wait until next year before buying life insurance
 - Our sense that these instances of practical reasoning are bad has nothing to do with assuming that certain counterpossibilities are salient to the reasoner → instead it is because the premise is not known
 - Basically, we can state it as follows: One does not know that p if they cannot use p for practical reasoning
 - Why? Because they do not know that p → isn't this circular
 - Might conflict with the epistemic possibility constraint
 - We can use p even when there is a non-zero (but super small) chance that not-p in practical reasoning → if there is a small epistemic chance that not-p, one knows p only if one is in a practical environment where the difference between a small chance that not-p and zero epistemic chance that not-p is irrelevant to matters at hand
- Advantage over contextualism
 - Suppose S uses 'I will be in Blackpool next year' as a premise in illegitimate practical reasoning. An ascriber, unaware of practical reasoning to which the premise is being put, asserts 'S knows that he will be in Blackpool next year'
 - By standard contextualist lights, that ascription is true → practical environment of the subject is not given any constitutive role to play in the truth conditions for knowledge ascriptions → connection between good practical reasoning and truth of knowledge ascriptions has been lost
 - Considerations of practical environment provide one clue as to why we may be overly pessimistic in our knowledge ascriptions: when we claim that no one can know that he or she will lose the lottery, part of what is going on is

that we realise that no one is in a position, in advance of the lottery draw, to acceptably sell a lottery ticket for minimal return

Stanley's Interest-Relative Invariantism

Interest-Relative Invariantism

- IRI: whether or not someone knows that p may be determined in part by practical facts about the subject's environment
 - Different variants such as probabilistic strength of evidence IRI: practical facts about a subject's environment at time t might make it the case that the subject must have stronger evidence than usual in order to know a proposition p at that time than she must possess in order to know that propositions at other times, where strength of evidence is measured in probabilistic terms
 - Relevant alternatives IRI is that the subject must rule out more or less and a different set of alternative propositions depending on practical facts regarding the environment
 - Practical facts are basically things that affect someone's practical reasoning and decision making
- Knowledge ascription is no longer relative to the context of the use of the word "knows", and has a fixed meaning regardless of who utters the sentence → the word 'know' expresses the same relation in every context
 - It is only relative to practical facts
- IRI is fundamentally the claim that knowledge is sensitive to an additional traditional non-epistemic factor i.e. practical facts about the environment
- Mechanism for why a proposition may be legitimately ignored
 - Small objective probability of it occurring, given my evidence
 - There may be nothing I can do about whether or not it will occur
 - The actions I would take if I knew it would occur would not result in a substantial improvement in my life over the plan of action I already intend to pursue
- It may be that a certain proposition is not a serious practical question for someone, but a consequence he or she may draw from it is a serious practical question and it may turn out that the negation of the consequence the agent draws has a sufficiently high probability, given the agent's interest and epistemic situation, to undermine the agent's knowledge of the consequence. Then the consequence is not known by that agent, even though it has been competently deduced by the agent from something known
 - Might licence the failure of intuitive closure: but Stanley says it is not obvious – for suppose that p is not a serious practical question for x at t. But then at t', x infers a consequence q from p, and q is a serious practical question for x at t'. Realising that q follows from p would change x's goals and interests. If so,

then p becomes a serious practical question for x at t' . This may result in undermining x 's knowledge of p at t' . So completely inferring a conclusion is in fact a way not only of gaining but also of undermining knowledge

- But I do not think this is so unintuitive → quite consistent with the characterisation of IRI
- Stanley characterises the notion of serious practical questions in terms of warranted expected utility: a proposition p is practically irrelevant iff where $a_1, \dots, a_n$ are actions at my disposal, the differences between the warranted expected utilities of $a_1, \dots, a_n$ relative to the nearest states of the world in which p are not meaningfully different from the differences between the warranted expected utilities of a_1 to a_n relative to the nearest states of the world in which not- p
 - So not just an agent's subjective credences
 - A proposition p is practically relevant iff its truth or falsity would affect the preference ordering of the actions at my disposal
- Appeal to serious epistemic possibility in addition to notion of serious practical question
 - Consider the proposition that you have an even number of hairs – I do not care how many hairs you have but I believe that you have an even number of hairs
 - If it turned out that you had an odd number of hairs, it would not matter – the various options I have at my disposal are to retain my belief that you would have an even number of hairs or give it up
 - Given that I do not care about the number of hairs you have, whether or not you have an odd number of hairs will not make a difference to the warranted expected utilities of retaining or discarding my belief – so the proposition that you have an odd number of hairs is not a serious practical question for me
 - Nevertheless, it is epistemically relevant to my belief that you have an even number of hairs, because it is a serious possibility in the epistemic sense that you have an odd number of hairs
 - The fact that the negation of a proposition is an epistemic possibility for an agent in a situation prevents the subject from knowing that proposition in question
- These serious epistemic possibilities and practical questions only undermine one's knowledge of a proposition p if, given her evidence, the probability of the negation of that proposition is not sufficiently low
- Example:
 - Low Stakes: the proposition that the bank is open on a Saturday is not a serious practical question for Hannah, so Hannah's evidence is sufficient to justify her belief that she knows that the bank will be open on Saturday

- High Stakes: the proposition that the bank is open on Saturday is a serious practical question for Hannah, since if it is not open, it would have a serious effect on the utility of some of her actions at her disposal
- Does not change even if the Attributor herself has high or low stakes
- Most difficulty with High Attributor-Low Subject stakes where the stakes for Hannah are high but the stakes for Bill are not High – predicts that Bill does know that the bank will be open on Saturday, though Hannah does not
 - Contradicts the intuition that Hannah's utterance of "Bill doesn't know that the bank will be open" expresses a true proposition
 - Hannah must say truly that "Bill does know that the bank will be open" even though she herself doesn't know → "S knows that P" entails P
- Stanley replies that this is not a prima facie concern – given intuitive connections between knowledge and practical reasoning, we should therefore be deeply suspicious of a theory that gives undue weight to our intuitions about High Attributor-Low Subject Stakes
- Comparison with Hawthorne's Subject-Sensitive Invariantism
 - Stanley says that SSI can account for Hannah's failure to assert that Bill knows that the bank will be open on Saturday (High Attributor, Low Subject) but not when High Attributor, High Stakes (asserting that Smith does not know)
 - Just because a proposition p is not assertible does not mean that its negation is assertible – if I have no good evidence for or against p, p is not assertible, but neither is not-p
 - Hawthorne wrongly explains that why Hannah can explain why the negation of the knowledge attribution is assertible is that to assert the knowledge attribution 'Bill knows that the bank will be open on Saturday' would implicate something false – namely, that Hannah reasonably takes herself to know that the bank will be open on Saturday – but it is not generally the case that 'not-S' is assertible if asserting S would implicate something false → does not rescue his theory
 - The explanation for the High Attributor, Low Subject case is that from a third person perspective, when we are asked for our intuitions about the case, we intuitively recognise that what Hannah and Sarah really care about is whether Bill would know, were he in Hannah and Sarah's practical situation → the answer to this is false
 - So Hannah would assert that Bill does not know
 - Differs from Hawthorne's explanation of overprojecting our lack of knowledge onto others
 - DeRose argues that if Hannah had lots of evidence relative to Bill, then it would say that she knows that the bank is open while Bill doesn't

know → and that this strategy cannot accommodate this data since Hannah is not projecting her lack of knowledge onto Bill, since by assumption, she knows that the bank is open

Advantage over Contextualism

- Problem with IRI: treatment of knowledge ascriptions in the scope of modal and temporal operators
 - According to IRI, for any p, knowing that p is a property, possession of which at a time depends in part upon how much is at stake for that person at that time
 - Suppose that Hannah knows at t that the bank will be open on Saturday, on the basis of having been there on a Saturday two weeks earlier; clearly, nothing much is then at stake for Hannah in the matter, but IRI will licence the truth of the following counterfactual
 - If Hannah had had a bill coming due, then she wouldn't know that the bank would have been open on Saturday
 - Past-tense: Hannah didn't know on Thursday that the bank would be open on Saturday, but she did know on Friday (she had a bill coming on Thursday but the company decided to cancel the bill on Friday)
- Contextualism is not burdened with the modal and temporal objections since it is independent of what is at stake for the person and instead affected by the context of the speaker
 - Stanley suggests that there might still be problems for contextualism: Hannah doesn't know that the bank will be open on Saturday. But if I wasn't aware that she had a bill coming due, she would know that the bank will be open on Saturday
 - If the property that is the 'epistemic' domain for 'know that the bank will be open on Saturday' is the property of being a possibility salient to the speaker, then, relative to a situation in which the speaker was not aware of the possibility that she has a bill coming due, this possibility would not be salient to the speaker – so relative to such a situation, what is actually expressed by "Hannah knows that the bank will be open on Saturday"
 - Must think of epistemic domains on the model of quantifier domains
- Problem of IRI and High-Attributor, Low Subject Stakes
 - Stanley: the interest-relative invariantist must deny that our intuitions about such cases as High Attributor-Low Subject Stakes are to be taken at face value
- Contextualism faces parallel issues
 - It is Friday, and Bill is considering going to the bank. Nothing much hangs on whether he gets to the bank before Monday. So he considers delaying going to the bank until Saturday. He has been to the bank before on Saturday, and found it open. So he assumes it will be open the next day (and it will). He

utters absent-mindedly, 'Well, since I know the bank will be open tomorrow, I'll just go tomorrow.' In contrast, it is very important for Hannah to get to the bank before Monday. She overhears Bill, and asks him what the basis for his claim was. But she too has been to the bank on previous Saturdays, which does not allay her concern that the bank has changed its hours. She therefore concludes that what Bill said was false.

- Intuitively, Hannah is right to conclude that Bill's utterance was false: but contextualism predicts that she is wrong → since nothing is at stake for Bill if he does not make it to the bank until Monday, he is in a low standards context and he does have sufficient justification to stand in the relation expressed by 'know', relative to that context, to the proposition that the bank will be open
- So both fail at the High Attributor, Low Subject stakes → for every such case, there is a precisely analogous case where the contextualist predicts the similar verdict as IRI
- Contextualists face problems with Ignorant High Stakes: Hannah does not know the cost of her actions → contextualists predict that she expresses a truth when she says she knows the bank will be open but intuitively, Hannah does not know
 - IRI provides a perfectly correct account
 - Contextualists can abandon the intention-based view of context-sensitive expressions: it is the actual interests of the participants in the conversational context that determine the semantic value of 'know', relative to a context, rather than what the participants perceive their interests to be: since it is in Hannah's interests (whether she knows it or not) to make it to the bank as soon as possible, she counts as being in a high stakes context
 - A similar manoeuvre allows the IRI-ist to account for High Attributor-Low Subject Stakes
 - The reason that IRI cannot provide a straightforward account is because it does not matter to Bill whether or not the bank will be open on Saturday
 - But if we expand the scope of the subject's interests at a time to include Hannah's interests among his interests → this is as unintuitive, Stanley claims, as abandoning the intention-based view of context-sensitivity
- Advantage of no semantic burden of justifying the contextualist claim that knowledge attributions are context-sensitive in a distinctively epistemological way
 - Firstly, IRI has no semantic burden to discharge since it is not a semantic thesis but a metaphysical thesis about the nature of the knowledge relation
 - Secondly, IRI gives a superior account of these cases, so the contextualist's evidential case is gone

- But the IRI must explain why, in an ordinary case, a skeptical scenario is not a serious practical alternative to one's ordinary beliefs
 - Stanley thinks this is difficult: altruism is much less meaningful when I am a BIV
 - But we can do so by assessing the probability that the event happens, which is very small
 - But this still doesn't deal with the skeptic that says the fact we have hands is not at all evidence that we are not BIVs → so we cannot objectively assess probabilities based on evidence
 - The IRI can mimic Lewis' notion of proper ignoring by changing the person who can properly ignore from the ascriber to the subject

Williamson's insensitive invariantism

Central Thesis: denies our intuitions that the semantic value of knowledge ascriptions change due to shifts in features of our practical environment or the context of ascription by providing a psychological explanation via the availability heuristic

- Contextualists and SSI attempt to provide a more charitable reference for 'know'
 - But Williamson argues that given reasonable assumption, any case of the kind that motivates claims of shifting standards is also liable to manifest cross-contextual patterns of dispositions to make mutually irreconcilable claims
 - Consider:
 - John: I know that this is a Zebra
 - Mary: How do you know that it isn't a mule cleverly painted to look like a zebra?
 - John: Hmmm, for all I know, it is a painted mule. So I was wrong. I don't know that it is a zebra
 - If SSI and Contextualists are right, then we cannot make sense of John's use of the past tense "I didn't know that it was a zebra after all" and that "I was wrong"
 - The dialogue is natural
 - The limits of assigning charitable reference is much more extensive than sensitive invariantists and contextualists expect
- The insensitive invariantists explanation via psychological biases
 - Apparent cases of standard shifting typically involve everyday ascriptions in low standards contexts
 - The insensitive invariantist is committed to the assertion of knowledge in low standard contexts being true and the denial of knowledge in high standard contexts being false: the anti-sceptical insensitive invariantist must explain the illusion of ignorance

- Availability heuristic: no choice but to judge knowledge ascriptions on an impressionistic basis, using convenient but highly fallible rules of thumb
 - We give things more weight when they are psychologically salient to us → ascriber gives newly mentioned evidence more weight
- So the sensitive invariantists and the contextualists lose their evidential advantage against insensitive invariantists
- Link between practical reasoning and knowledge does not favour all applications of contextualism and sensitive invariantism
 - KPR: A first-person present-tense ascription of 'know' with respect to a proposition is true in a context iff that proposition is an appropriate premise for practical reasoning in that context
 - So for a proposition p that makes little practical difference and p^* that makes a huge practical difference we have:
 - $K[S, p, t, w]$ is true in C iff p is an appropriate premise for practical reasoning in C (a)
 - $\text{not-}K^*[S^*, p^*, t^*, w^*]$ is true in C^* iff p^* is not an appropriate premise for practical reasoning in C^* (b)
 - Appeal of KPR is there but it is not obviously true
 - Cases that may not work well for the contextualist/sensitive invariantist
 - Consider an epistemology seminar with a focus on radical skeptical and in the middle of the seminar, a friend offers you a bet in which you gain \$5 if you have hands and lose \$50 otherwise → you will take the \$5 (ignoring the fact that the skeptical scenario raises problems for enforcing the bet) → conflicts with KPR
 - KPR for the insensitive invariantist
 - KPR+ (simplified): One knows q iff q is an appropriate premise for one's practical reasoning → so when q is appropriate but one is in no position to know that q is appropriate, one in effect knows q without being in a position to know that one knows q
 - Anti-luminosity: one can know q without being in a position to know that one knows q
 - The insensitive invariantists can explain why (a) and (b) look jointly plausible even though the agent in C^* is epistemically as well positioned with respect to p^* as the agent in C is with respect to p
 - For suppose that the agent in C is in no position to know that she knows p and the agent in C^* is in no position to know that she knows p^*
 - Since the stakes are higher in C^* than in C , we may view the failure of Second Order knowledge in C^* more sternly than its failure in C , and hence regard p as appropriate in C but not in C^*

Nagel's Invariantism

Knowledge Ascriptions and Thinking about Error

Central Thesis: Williamson and Hawthorne explain the increased stringency of intuitive ascriptions of knowledge when unrealised possibilities of error are mentioned as due to our wrong overestimation of error due to a heuristic bias. Nagel denies this and argues that the temptation is instead best explained via epistemic egocentrism and hindsight bias

- Problem Cases we are considering
 - Plain Story A: John Doe is in a furniture store. He is looking at a bright red table under normal lighting conditions. He believes the table is red. Does he know that the table is red
 - Detailed Story B: John Doe is in a furniture store and is looking at a bright red table under normal lighting conditions. He believes that the table is red, but a white table under red lighting would look exactly the same to him and he has not checked whether the lighting is normal, or whether there might be a red spotlight shining on the table. Does he know that the table is red.
 - Anti-skeptical invariantists such as Williamson (and Hawthorne since nothing about the subject's context has changed) would want to say that John knows that the table is red in both cases A and B
 - Contend that 'knows' always picks out the same relation
- Williamson's Availability Heuristic
 - The fundamental point is that Williamson argues that our intuitions are mistaken: we err in retracting knowledge about third parties in cases B when unrealised possibilities about error are made salient to us
 - Williamson and Hawthorne both invoke Tversky and Kahneman's availability heuristic in which we overestimate the likelihood of an event when we can easily remember or imagine it, resulting in inaccurate judgements → for instance when asked about frequency of various causes of deaths, people overestimate the deaths by tornado etc
 - Hawthorne: "When certain non-knowledge-destroying counterpossibilities are made salient, we overestimate their real danger; as a result, we may find ourselves inclined to deny knowledge to others in cases where there is in fact no real danger of error"
- Problems with the availability heuristic
 - Empirics
 - Problem 1: Easy to imagine cases do not actually increase our estimation of likelihood in all cases and instead sometimes underestimate them
 - For example, Oppenheimer's work on the spontaneous discounting of availability: people tend to underestimate the

frequency of surnames such as “Bush”, “Clinton” – natural psychological inclinations to single rather than multiple causes

- The invariantist is too quick to say that the increased availability of unrealised error immediately increases the apparent probability of error
- Automatic discounting of significance of availability whenever an offsetting explanation of increased availability is present

- Problem 2: To say that a heuristic principle can generate a judgement error in some circumstances is not to say that it would do so across the board – invariantists who are relying on the availability explanation of knowledge ascriptions shifts the need to defend the thesis that this effect is operative in contexts in which we are attempting to engage in careful reasoning about epistemological issues

- Such a thesis runs contrary to a more general psychological fact about accountability: reliance on most heuristics tend to be sharply attenuated or cancelled in conditions in which we have an incentive for accuracy

- Conceptual

- Problem 1: Not clear that odds of danger make the right sort of difference to knowledge ascriptions – even if we face an unrealised possibility of error, we would adjust it for its probability of occurrence → so it's not even the case that it may affect our knowledge ascriptions

- But the whole point is that it is irrational

- Problem 2: Under the working hypothesis that knowledge ascriptions always involve some assessment of the level of epistemic danger, if one concedes that sceptical scenarios might be so easy to imagine that frequent rehearsal would increase the odds of epistemic danger, then even plain cases should strike typical epistemologists as cases of failure to know

- Even if discounting kicks in and the odds of danger appear lower than they should when error is mentioned in B, the fact that they appear at all could explain our tendency to refrain from ascribing knowledge
- What would matter on this suggestion would not be quantitative contrasts in the apparent odds of danger but qualitative contrasts between salience and non-salience of various prospects of error

- Nagel's epistemic egocentrism and hindsight bias

- Nagel supports Williamson's suggestion that knowledge is a mental state and agrees that we do make some mistake in assessing case (b) because we superimpose our own third-party mental condition as the subject's mental condition
 - Epistemic egocentrism
- Epistemic egocentrism takes the form of hindsight bias: our capacity to recall our own less informed past states of mind is tainted by more recently acquired information → affects our efforts to judge the mental states of less informed others
 - We overestimate the extent to which they share our beliefs, attitudes and concerns
- Application to case B: once concerns about the possibility tricky lighting have been raised for me, I illegitimate John's predicament as if he shared those concerns
 - Ordinarily, a person who is actively concerned about the lighting would glance up to check it prior to making a judgement about the colour of the table
 - Egocentrism drives me to evaluate John's cognition as if he shares all my concerns, he will look hasty to me and seem to fall short of knowledge, only because I am illegitimately importing a feature of my own condition

Knowledge Ascriptions and Changing the Stakes

Central Thesis: Nagel rejects both Contextualism and IRI in that knowledge is sensitive to not epistemic features and that the reason for the change in our intuitions for knowledge ascriptions when practical interests are mentioned is due to a need-for-closure and a change in one's confidence in their beliefs

- Epistemic Test Cases
 - Lo and her next-door neighbour Hi are thinking about going out to the bank, but wondering whether the trip could be postponed to tomorrow – both of them remember a recent Saturday visit to the bank but neither of them remember further information relevant to the bank's opening
 - For Lo, the question is of no importance because her errand can be done anytime; for Hi, the question has burning practical importance since Hi must deposit his paycheck before Monday, otherwise he will default on his mortgage and lose his home
 - Reasons for why we attribute more knowledge to Lo than Hi changes
 - Contextualist: shifts indicates that the referent of 'know' changes with the situation of the ascriber – when we read about Hi's predicament our epistemic standards become more demanding than they are when we read about Lo

- SSI or IRI: focus on who the knowledge is ascribed to – features of practical environment change
 - Invariantism: only truth conducive factors count; so typically they attempt to dismiss these intuitions of shift, construing them as the products of psychological bias or failures to observe the distinction between knowing and knowing that one knows
- Need-for-closure: the bank cases can be explained by a psychological effect often causally related to, but distinct from, perceived stakes – so the evidential value of competing accounts softens and goes to almost zero
 - Closure marks the switch from the formation to the possession of belief
 - After closure we have a subjective sense of a stable belief
 - Prior to closure we are searching for information and trying to figure out what to believe in and we are weighing various alternatives → not confident in our beliefs
 - Subjects with different levels of need-for-closure can differ in whether they have or lack knowledge for the straightforward reason that they have or lack a solid belief and we need not need to resort to non-epistemic factors to explain the shifts
 - High stakes subjects will search for more information before making up their minds and process this information more thoroughly, for example, by selecting more complex decision strategies and assigning different weights to more or less important attributes
 - Often also give consideration to less important factors by sheer availability (dilution effect) → what increases is more the cognition rather than the accuracy itself due to minor biases
- Against Contextualism and IRI: if certain stake-related psychological effects do explain the epistemic shifts in the practical interest cases used to support contextualism and IRI, these cases do not have the evidential value against invariantism that they claim to have → Nagel argues that the intellectualist invariantist will always have a psychologically plausible story about why the epistemic intuitions shift
 - Bach argues that what really changes between Low and High stakes is the subject's threshold for confident belief in the proposition that the bank is open on Saturday: in High, her true belief is insufficient confident to count as knowledge → have to construct a threshold of confident belief to count as belief
 - Whether one thinks of High-Stakes Hannah as having lower-confidence belief or a state of cognition that precedes arrival at a settled belief, it is psychologically realistic to read her as needing more evidence to either make up her mind at all or attain the same level of subjective confidence in High Stakes as she would have in Low

- High-Stakes Hannah is a need-for-closure subject: she is strongly averse to making a mistake about the baking hours and is willing to entertain hypotheses that makes the initial evidence inconclusive
 - Makes sense of Bach's objection: we ascribe knowledge in Low Stakes and deny in High stakes because we naturally attribute higher and lower confidence beliefs because we naturally attribute higher and lower confidence belief to the contrasted subjects, or confident belief and a state of evidence assessment that precedes the fixed belief
 - Because of the perceived need for closure differences between the subjects, we don't expect the same information to produce the same level of belief in High and Low stakes
 - Not about interest itself but about need-for-closure and belief